



*High performance. Delivered.*

# **Parallel processing with Spring Batch**


## Lessons learned

# accenture<sup>></sup> Morten Andersen-Gott

*High performance. Delivered.*

- Manager at Accenture Norway
- 30 years old
- Been using Spring Batch since 1.0M2 (2007)
- Member of JavaZone program committee
  - <http://tinyurl.com/javaever>
  - <http://tinyurl.com/thestreaming>
  - <http://tinyurl.com/ladyjava>

 mortenag

 [www.github.com/magott](http://www.github.com/magott)

 [www.andersen-gott.com](http://www.andersen-gott.com)



- Functional background for the batch
- Short introduction to Spring Batch
- Even shorter on Hibernate
- The problems
- The problems
- The problems

# Norwegian Public Service Pension Fund (SPK)

- Norway's main provider of public occupational pensions
- Also provide housing loans and insurance schemes
- Membership of the Norwegian Public Service Pension Fund is obligatory for government employees
- Stats
  - 950,000 members across 1600 organisations
  - Approx 138,000 receive a retirement pension
  - Approx 58,000 receive a disability pension
  - 950,000 members have total accrued pension entitlements in the Norwegian Public Service Pension Fund of 339 billion kroner.

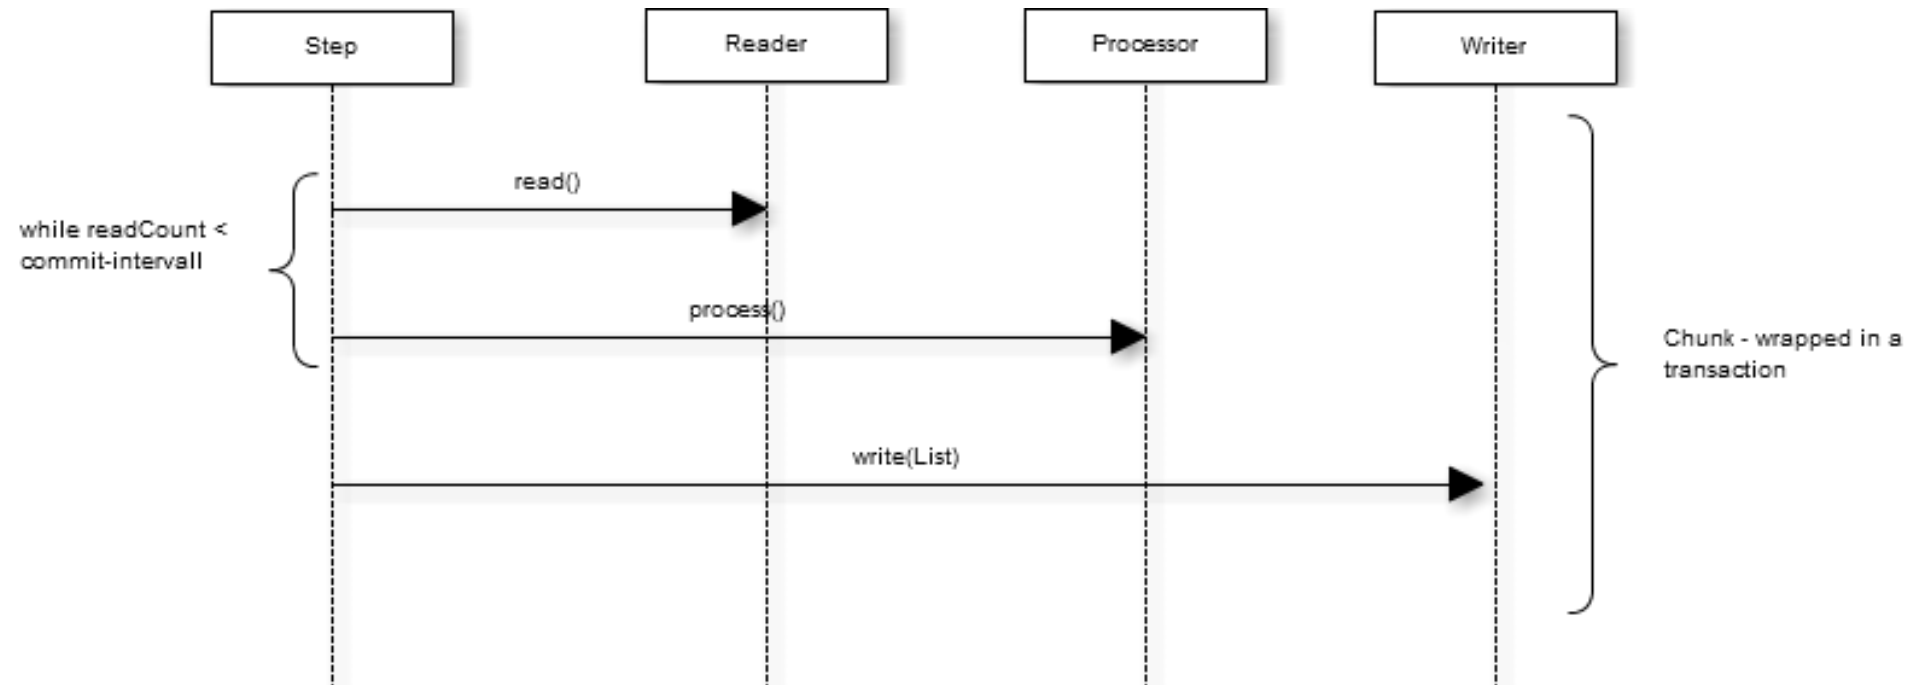
- Every year the parliament sets the basic amount of the national insurance
- This amount is a constant used in calculation of all benefits
- When the basic amount is changed, all benefits must be recalculated
- It's more complex than a constant in an algorithm
  - Rules are result of political games the last 50 years
  - Complex rules
  - Lots of exemptions

# Execution time requirements

- SPK's recalculation batch must run
  - After the basic amount is set
  - After the Labour and Welfare Administration has done it's calculation
  - Before the pensions are due next month
- Window of 1 week
  - Will ideally only run during weekends
  - Can not run while case workers are doing their job

- Framework for developing batch applications
- Implements batch semantics
  - Steps, Chunks, Stop, Restart, Skips, Retries
  - Partitioning, Multithreading, Parallel steps

# A step...





- **ItemReader**
  - `read()` returns one row at the time
  - Step is completed once `read()` returns `null`
- **ItemProcessor**
  - `process(item)` item is return value from `read()`
  - Business logic goes here
  - Items can be filtered out by returning `null`
- **ItemWriter**
  - `Write(list)` list of items returned from `process()`

```
<job id="foo">
  <step id="fileImport">
    <tasklet>
      <chunk commit-interval="10"
        reader="reader"
        processor="processor"
        writer="writer"/>
    </tasklet>
  </step>
</job>
<bean id="reader" class="...">
<bean id="processor" class="...">
<bean id="writer" class="...">
```

- A chunk is a unit of work
  - Executes within a transaction
  - Size is defined by number of items read
- A step is divided into chunks by the framework
- When  $x$  is the chunk size
  - `read()` is called  $x$  times
  - `process()` is called  $x$  times
  - `write` is called once with a list where `list.size() == x` (minus filtered items)

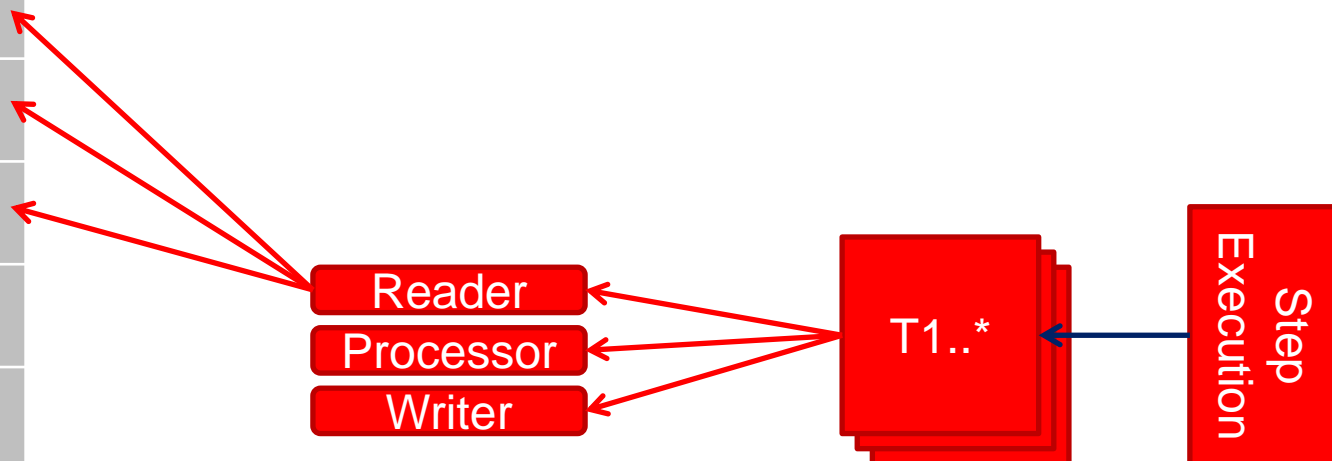


*High performance. Delivered.*

# Scaling

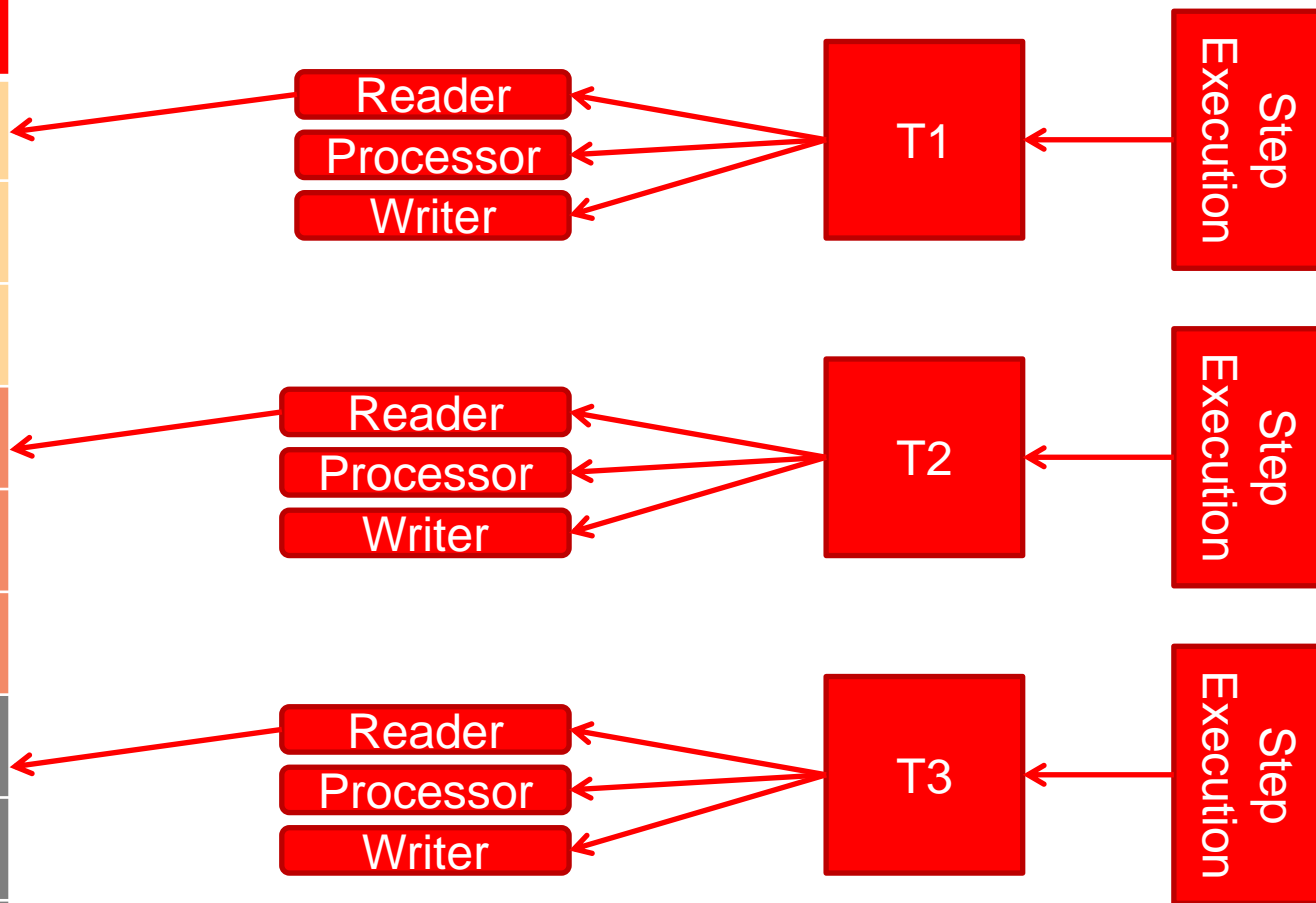
# Multi-threaded step

Id	Name
1	Paul
2	John
3	Lisa
4	Simon
5	Rick
6	Julia
7	Olivia
8	Scott
9	Martin



# Partitioned Step

Id	Name
1	Paul
2	John
3	Lisa
4	Simon
5	Rick
6	Julia
7	Olivia
8	Scott
9	Martin



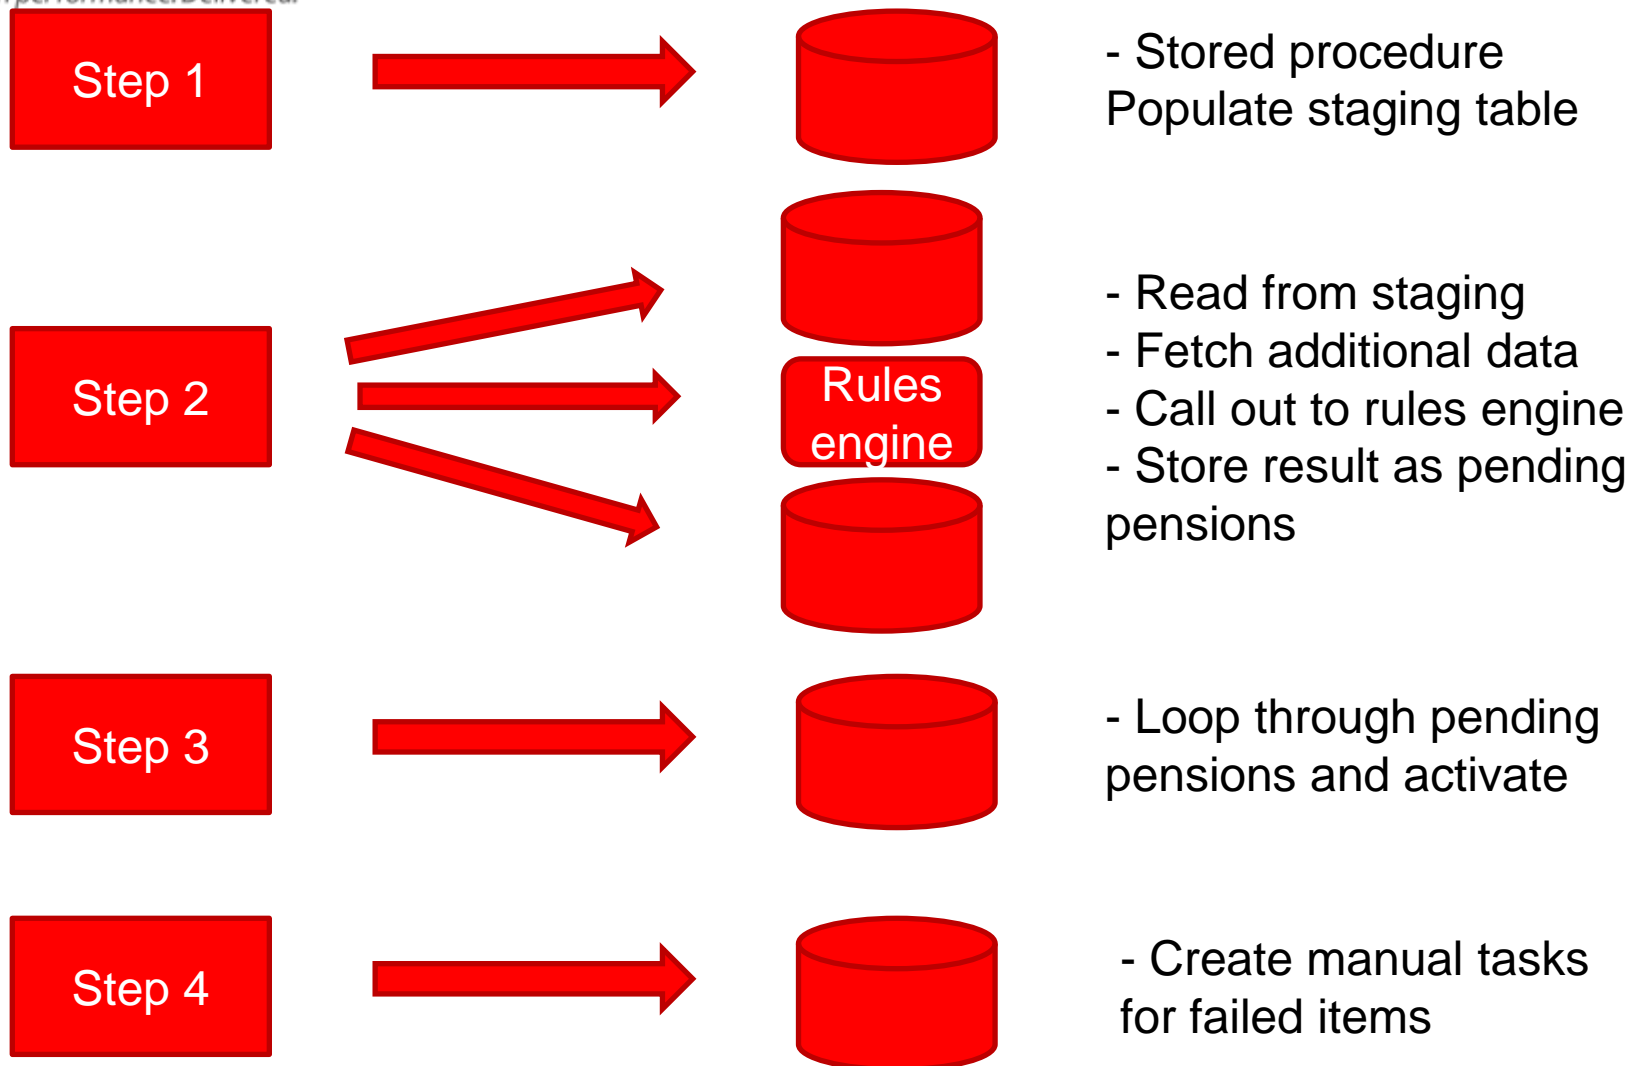
- Proxy
- Session cache
- Flushing
  - Queries
  - Commit



*High performance. Delivered.*

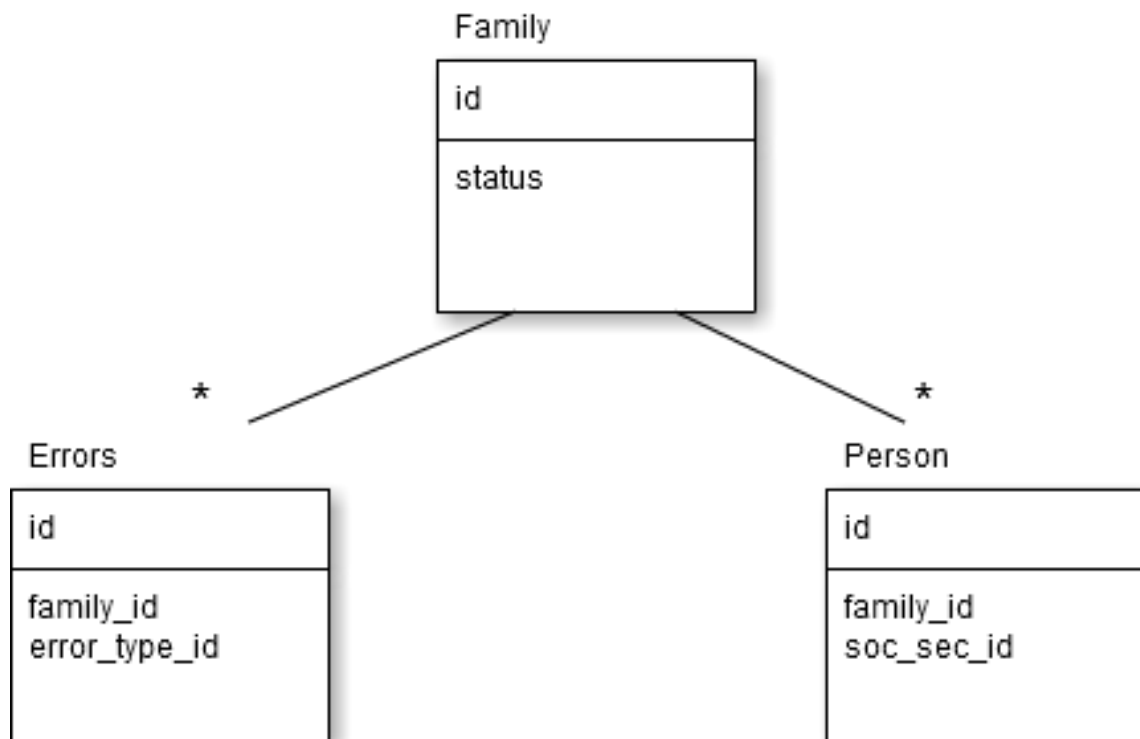
# **The pension recalculation batch**

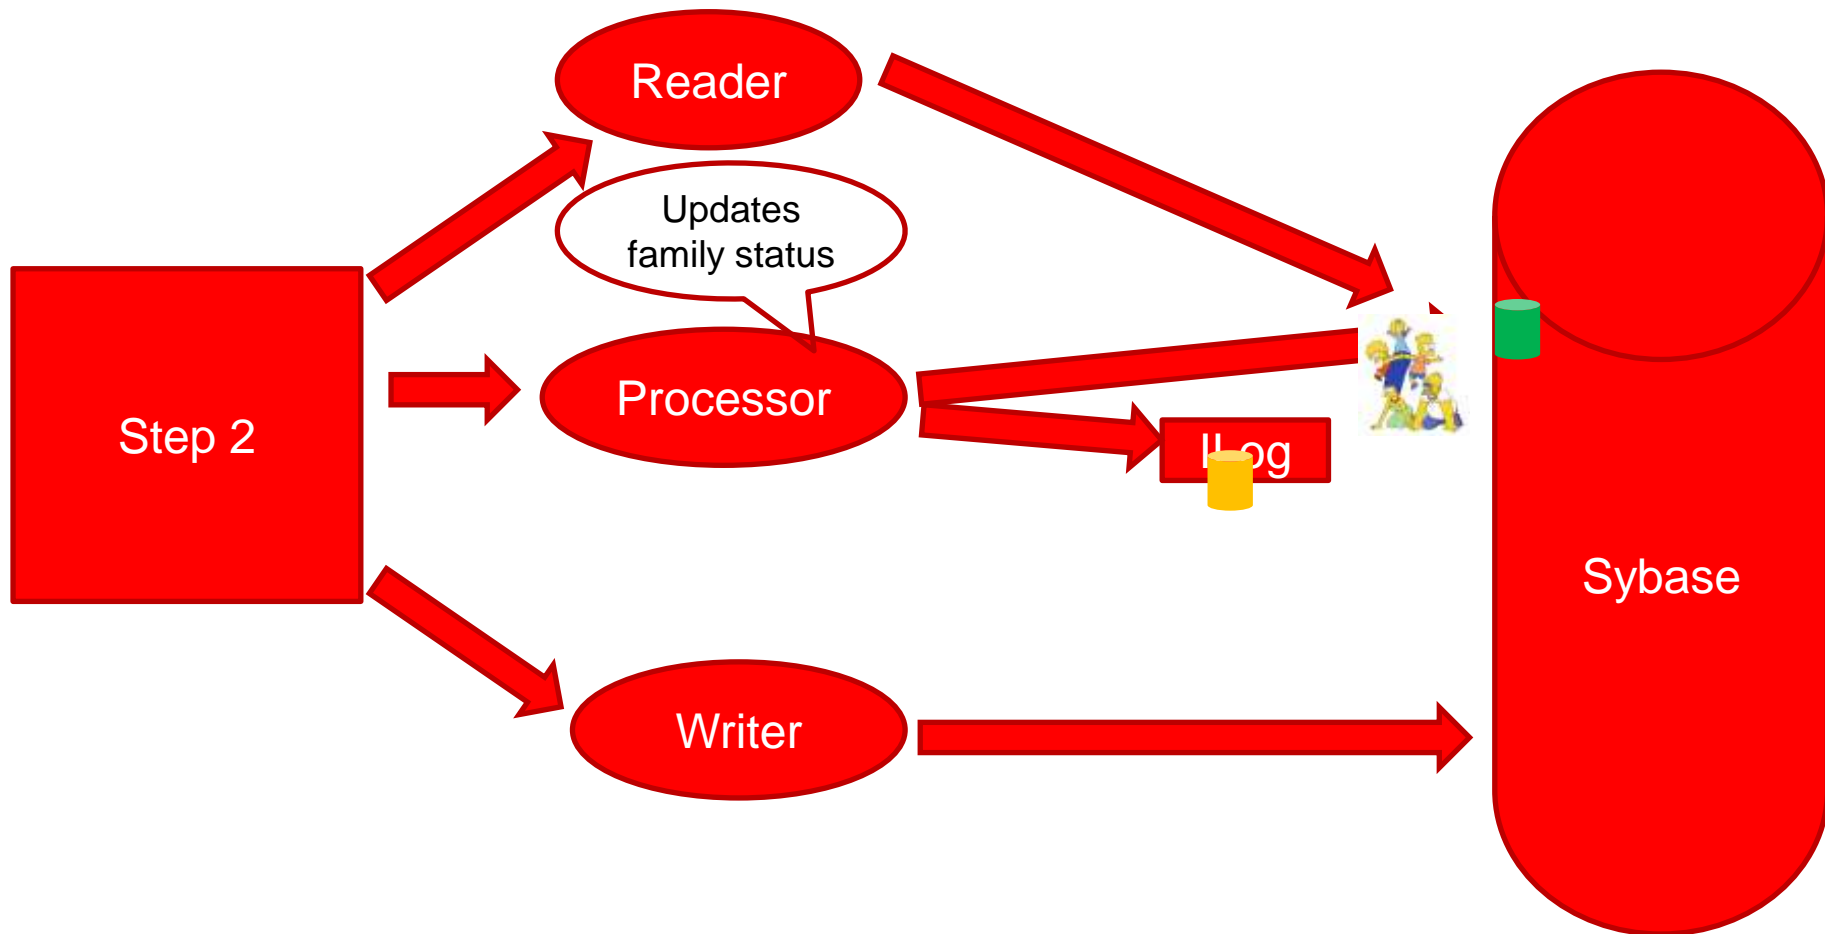




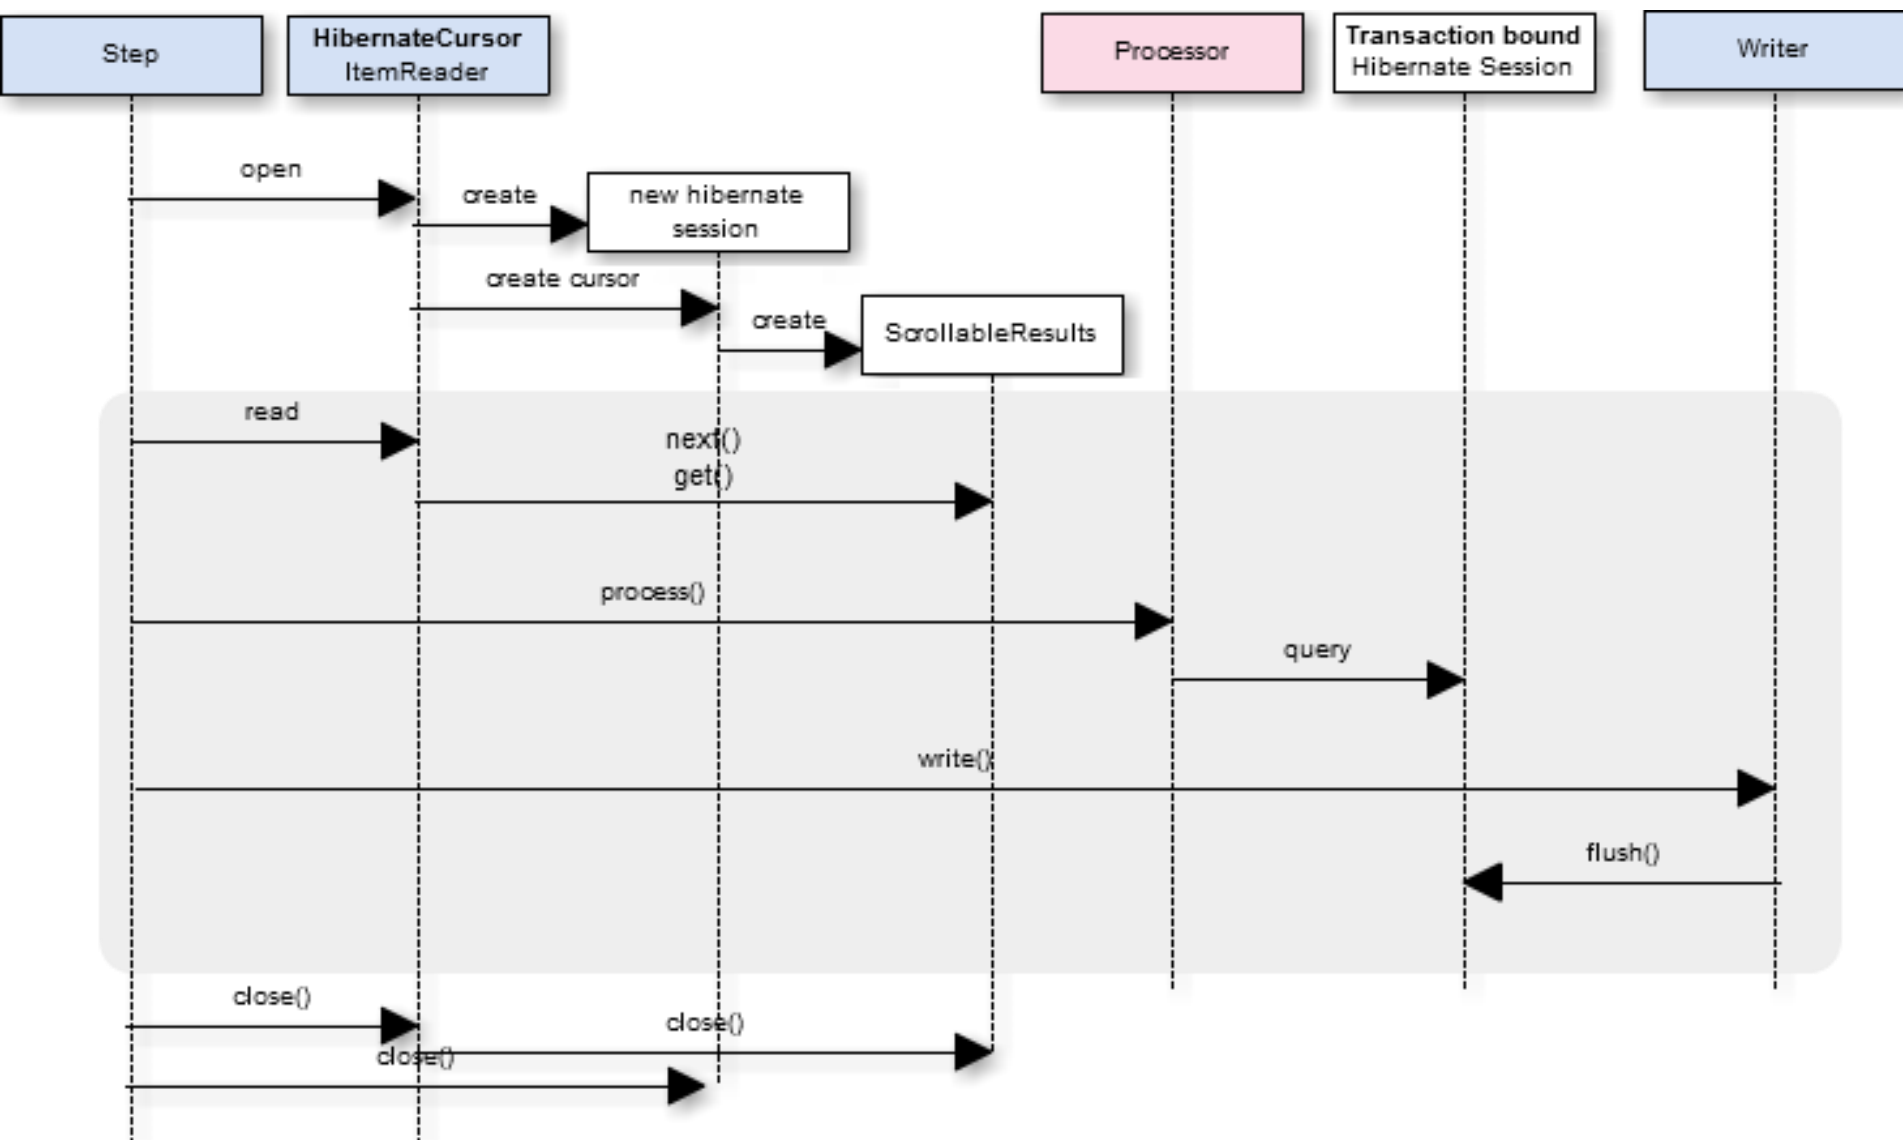
- A batch processing pattern
- A simple table with a identity column, functional key and processing status
- Restart is easy
  - Select soc\_sec from staging where processed=false
- An easy way to get parallel processing capabilities

# Our staging table(s)





- Using a separate statefull session in reader
  - Not bound to active transaction
  - Is never flushed
  - `clear()` is called before commit
- Entities are written to database using a (different) transaction bound session
  - Used by processor and writer





*High performance. Delivered.*

# Problems?

## What happened?

org.hibernate.HibernateException: Illegal  
attempt to associate a collection with two  
open sessions



- Family.errors and Family.persons are attached to reader's session
- Attempting to attach them to transaction bound session
- Hibernate will have non of that!



*High performance. Delivered.*

# The solution?

## 13.3. The `StatelessSession` interface

Alternatively, Hibernate provides a command-oriented API that can be used for streaming data to and from the database in the form of detached objects. A `StatelessSession` has no persistence context associated with it and does not provide many of the higher-level life cycle semantics. In particular, a stateless session does not implement a first-level cache nor interact with any second-level or query cache. It does not implement transactional write-behind or automatic dirty checking. Operations performed using a stateless session never cascade to associated instances. Collections are ignored by a stateless session. Operations performed via a stateless session bypass Hibernate's event model and interceptors. Due to the lack of a first-level cache, `Stateless` sessions are vulnerable to data aliasing effects. A stateless session is a lower-level abstraction that is much closer to the underlying JDBC.

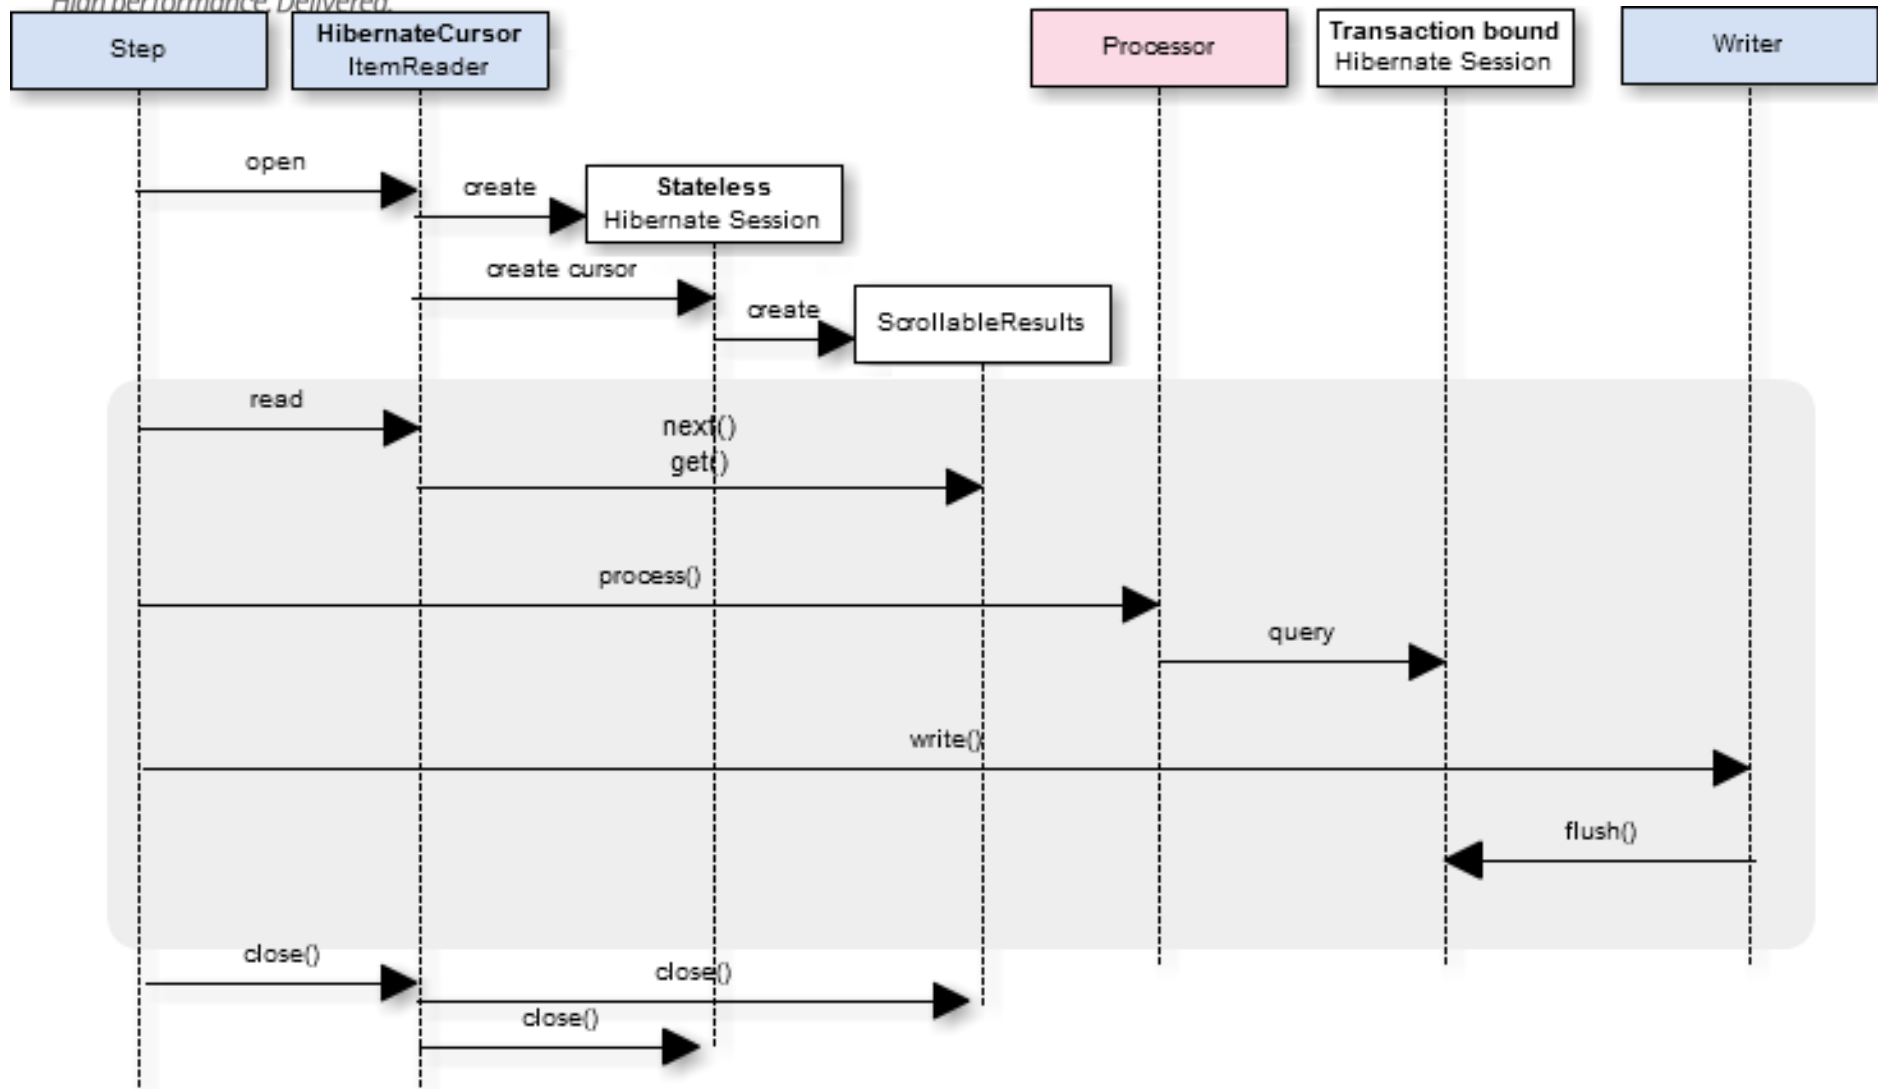
# Hibernate in the reader

(Second attempt)

- Let's try stateless session
- Default behavior in Spring Batch's Hibernate ItemReaders
  - useSateless="true"
- LEFT JOIN FETCH for eager loading collections
  - Avoiding LazyLoadingExceptions

# Hibernate in the reader (Second attempt)

High performance Delivered.





*High performance. Delivered.*

# Problems?




## What happened?

- `org.hibernate.HibernateException: cannot simultaneously fetch multiple bags`

- Hibernate is unable to resolve the Cartesian product
  - Throws exception to avoid duplicates



# You are using it wrong!

<p><b>max</b></p> <p>Hibernate Team</p>  <p><b>Joined:</b> Tue Aug 26, 2003 6:10 am  <b>Posts:</b> 8604  <b>Location:</b> Neuchatel, Switzerland (Danish)</p>	<p><b>Post subject:</b></p> <p><b>Posted:</b> Fri May 19, 2006 2:51 am</p> <p>It did <i>*not*</i> work in previous versions - it resulted in bags that could contain redundant elements because of the cartesian product.</p> <p>People did not realize that so now we are forcefully complaining when you specify that on queries/mappings.</p> <hr/> <p>Max</p> <p>Don't forget to rate</p>
<p>Top</p>	<p> <a href="#">profile</a></p>
<p><b>emmanuel</b></p> <p>Hibernate Team</p>  <p><b>Joined:</b> Sun Sep 14, 2003 3:54 am  <b>Posts:</b> 7173  <b>Location:</b> Atlanta, USA</p>	<p><b>Post subject:</b></p> <p><b>Posted:</b> Fri May 19, 2006 12:44 pm</p> <p>The main problem is that bag semantic (through Collection or List (wo @IndexColumn) ) is way overused.  95% of collections should really be a Set</p> <hr/> <p>Emmanuel</p> <p>Check Hibernate Search in Action out</p>

# Curing the *you are using it wrong* syndrome

# Hibernate in the reader

## (The third attempt)

- Examine the object graph
- Replace List with Set
  - Only one eagerly loaded collection may be of type list List
- This works...
  - ..for a while
  - We'll revisit Hibernate later...

- New demands sneak in
  - The batch should not abort
    - Not under any circumstance
- The batch should deal with
  - Missing or functionally corrupt data
  - Programming errors
  - ...

```
try{  
    //do data and business operations  
}catch(Exception e){  
    //Add error to staging & continue  
    family.addError(createError(e));  
}
```



*High performance. Delivered.*

# Problems?

# Overzealous exception handling

an assertion failure occurred (this may indicate a bug in Hibernate, but is more likely due to unsafe use of the session)

# Overzealous exception handling

## The solution

- Some exception **MUST** result in a rollback
  - Ex. `StaleStateException`
- Configure the framework to do retry/skip for these
- Only catch exceptions you **know** you can handle in a meaningful way
  - Nothing new here
  - Do not succumb to crazy requirements



# Time to turn on parallelization

- We chose partitioned over mutli-threaded step
  - No need for a thread safe reader
    - Step scope
    - Each partition get a new instance of reader
  - Page lock contentions are less likely
    - Row 1 in partition 1 not adjacent to row 1 in partition 2

- Legacy database using page locking
- Normalized database
- Relevant data for one person is spread across a number of tables
- Different threads **will** access same data pages
- Deadlocks **will** occur

# Page locking

ID	NAME	
1	Paul	T1
2	John	T2 waiting
3	Simon	
4	Scott	T1 waiting
5	Lisa	T2
6	Jack	
7	Nina	
8	Linda	T3



*High performance. Delivered.*

# DeadlockLoserDataAccessException

# Retry to the rescue

```
<step id="step2">
  <tasklet>
    <chunk reader="reader" processor="processor" writer="writer"
      commit-interval="10" retry-limit="10">
      <retryable-exception-classes>
        <include class="...DeadlockLoserDataAccessException"/>
      </retryable-exception-classes>
    </chunk>
  </tasklet>
</step>
```

# The mystery exception

```

/**
 * Thrown when a version number or timestamp check failed, indicating that the
 * <tt>Session</tt> contained stale data (when using long transactions
 * with versioning). Also occurs if we try delete or update a row that does
 * not exist.<br>
 * <br>
 * Note that this exception often indicates that the user failed to specify the
 * correct <tt>unsaved-value</tt> strategy for a class!
 *
 * @see StaleObjectStateException
 * @author Gavin King
 */
public class StaleStateException extends HibernateException {





    public StaleStateException(String s) {
        super(s);
    }
}

```

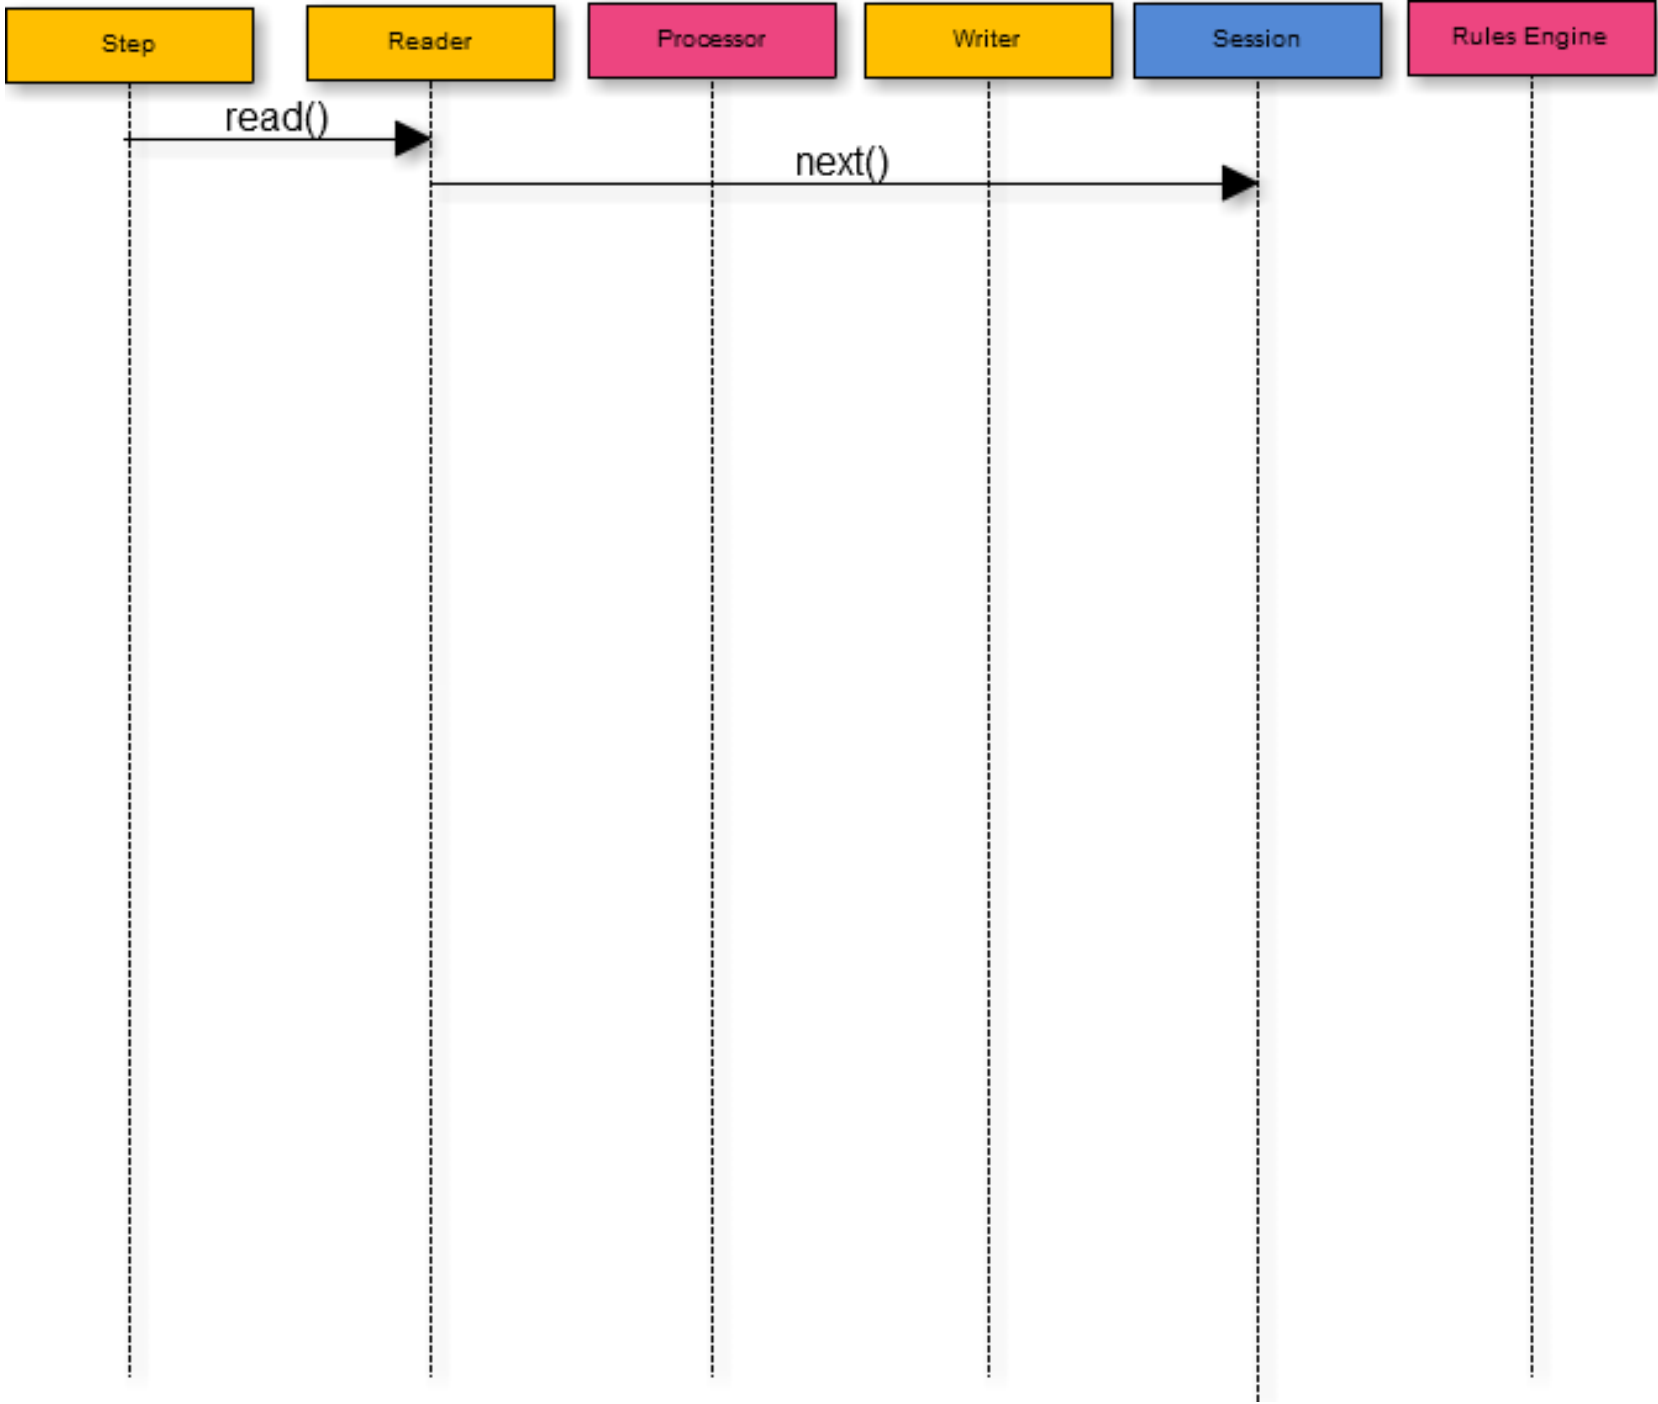


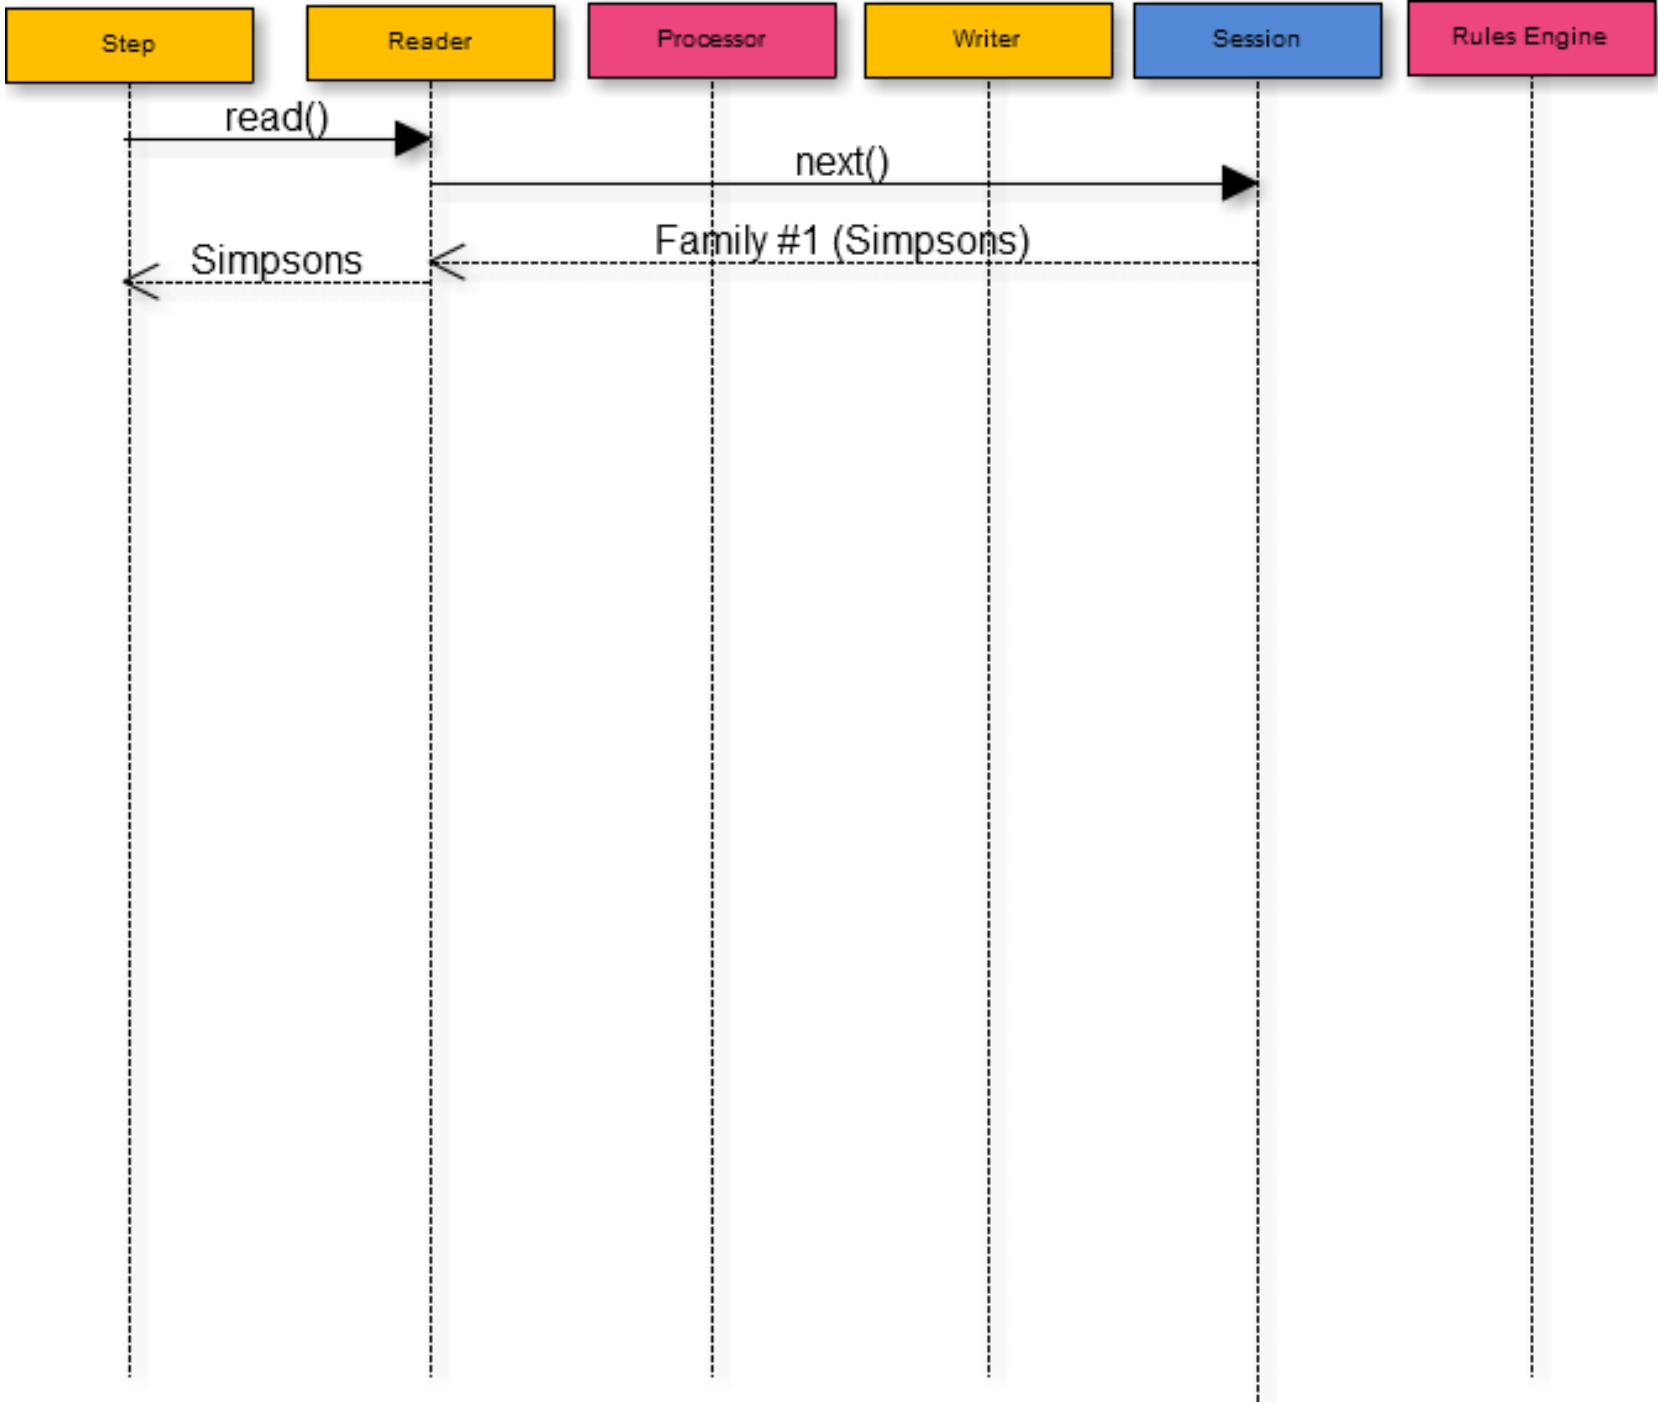
*High performance. Delivered.*

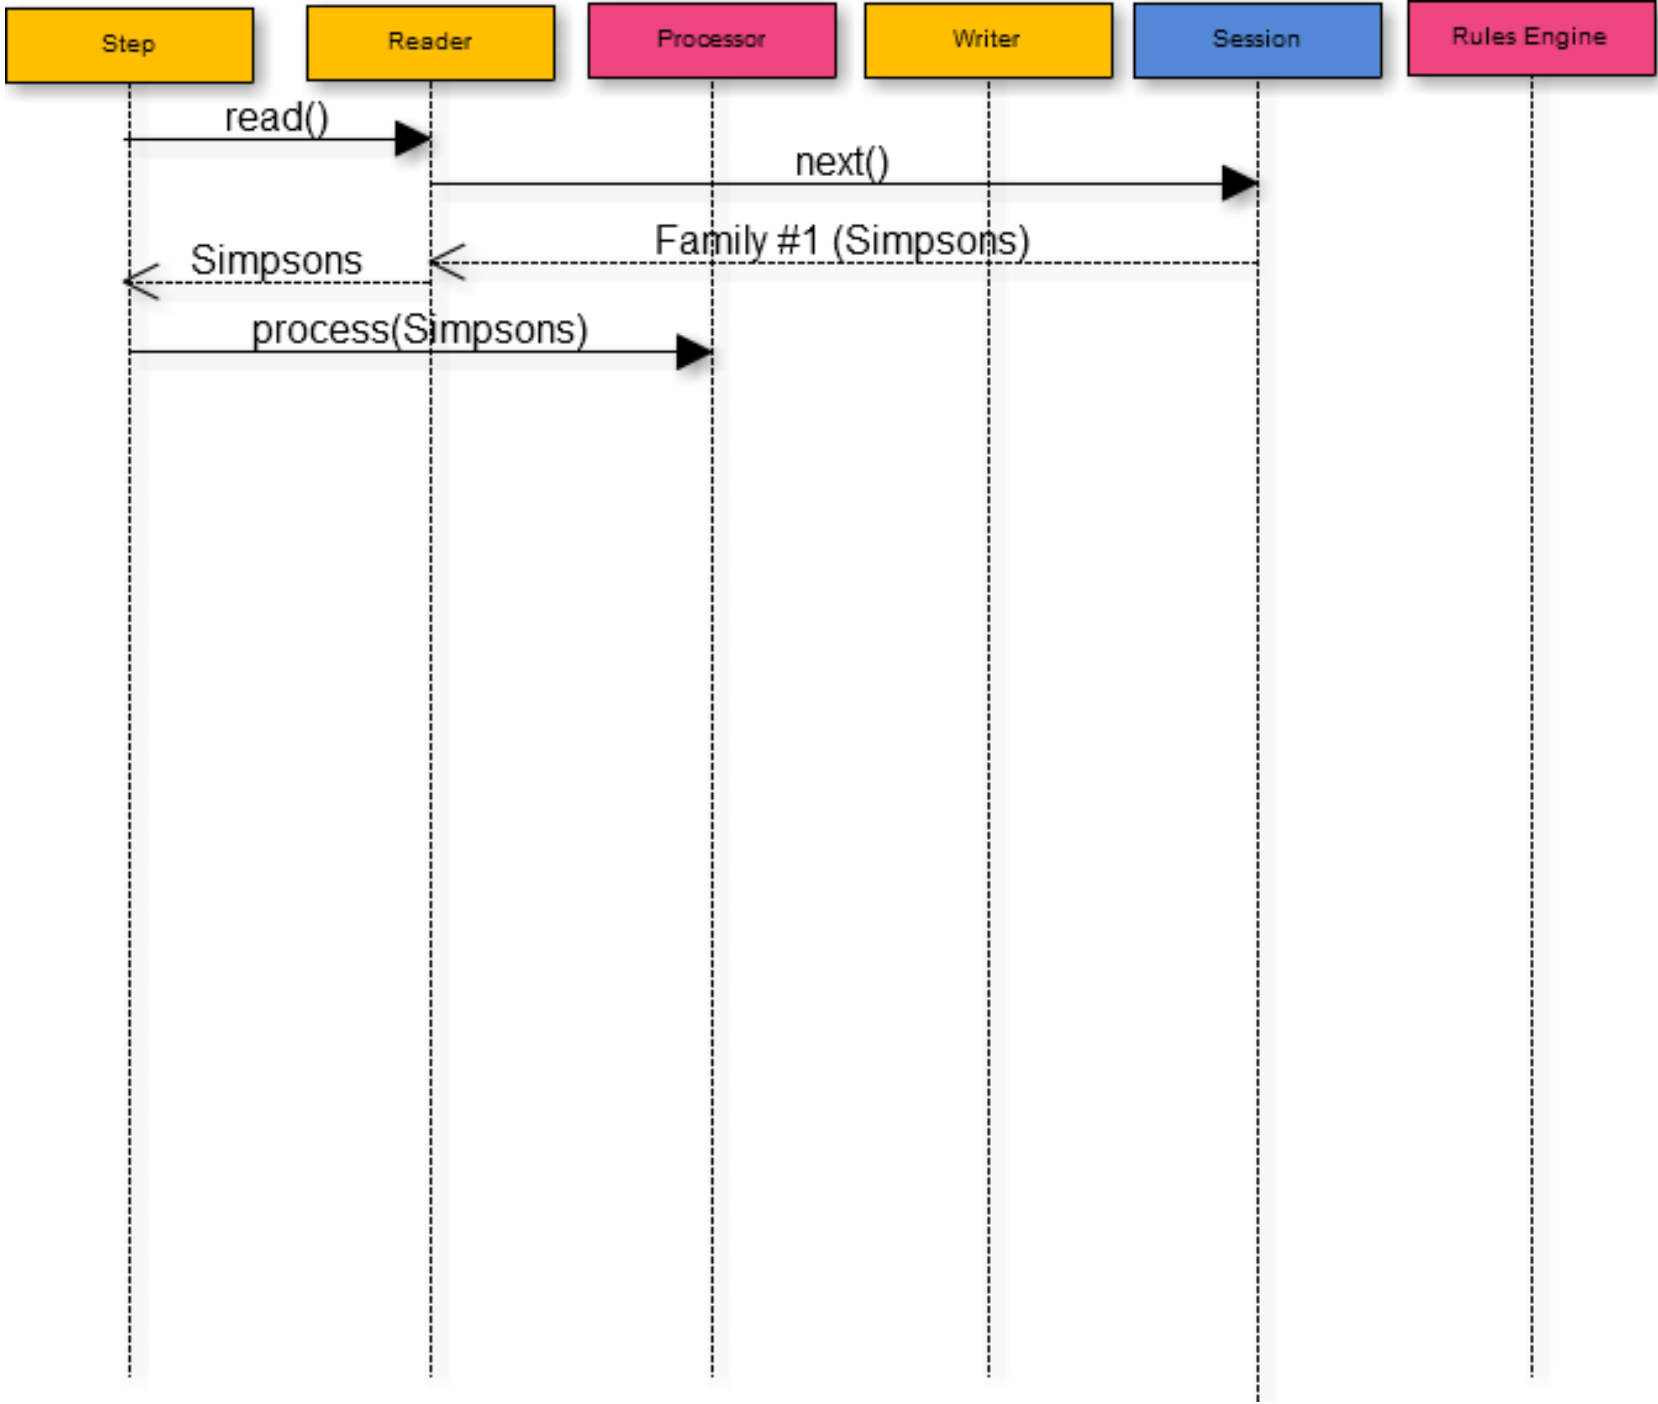
Two weeks later..

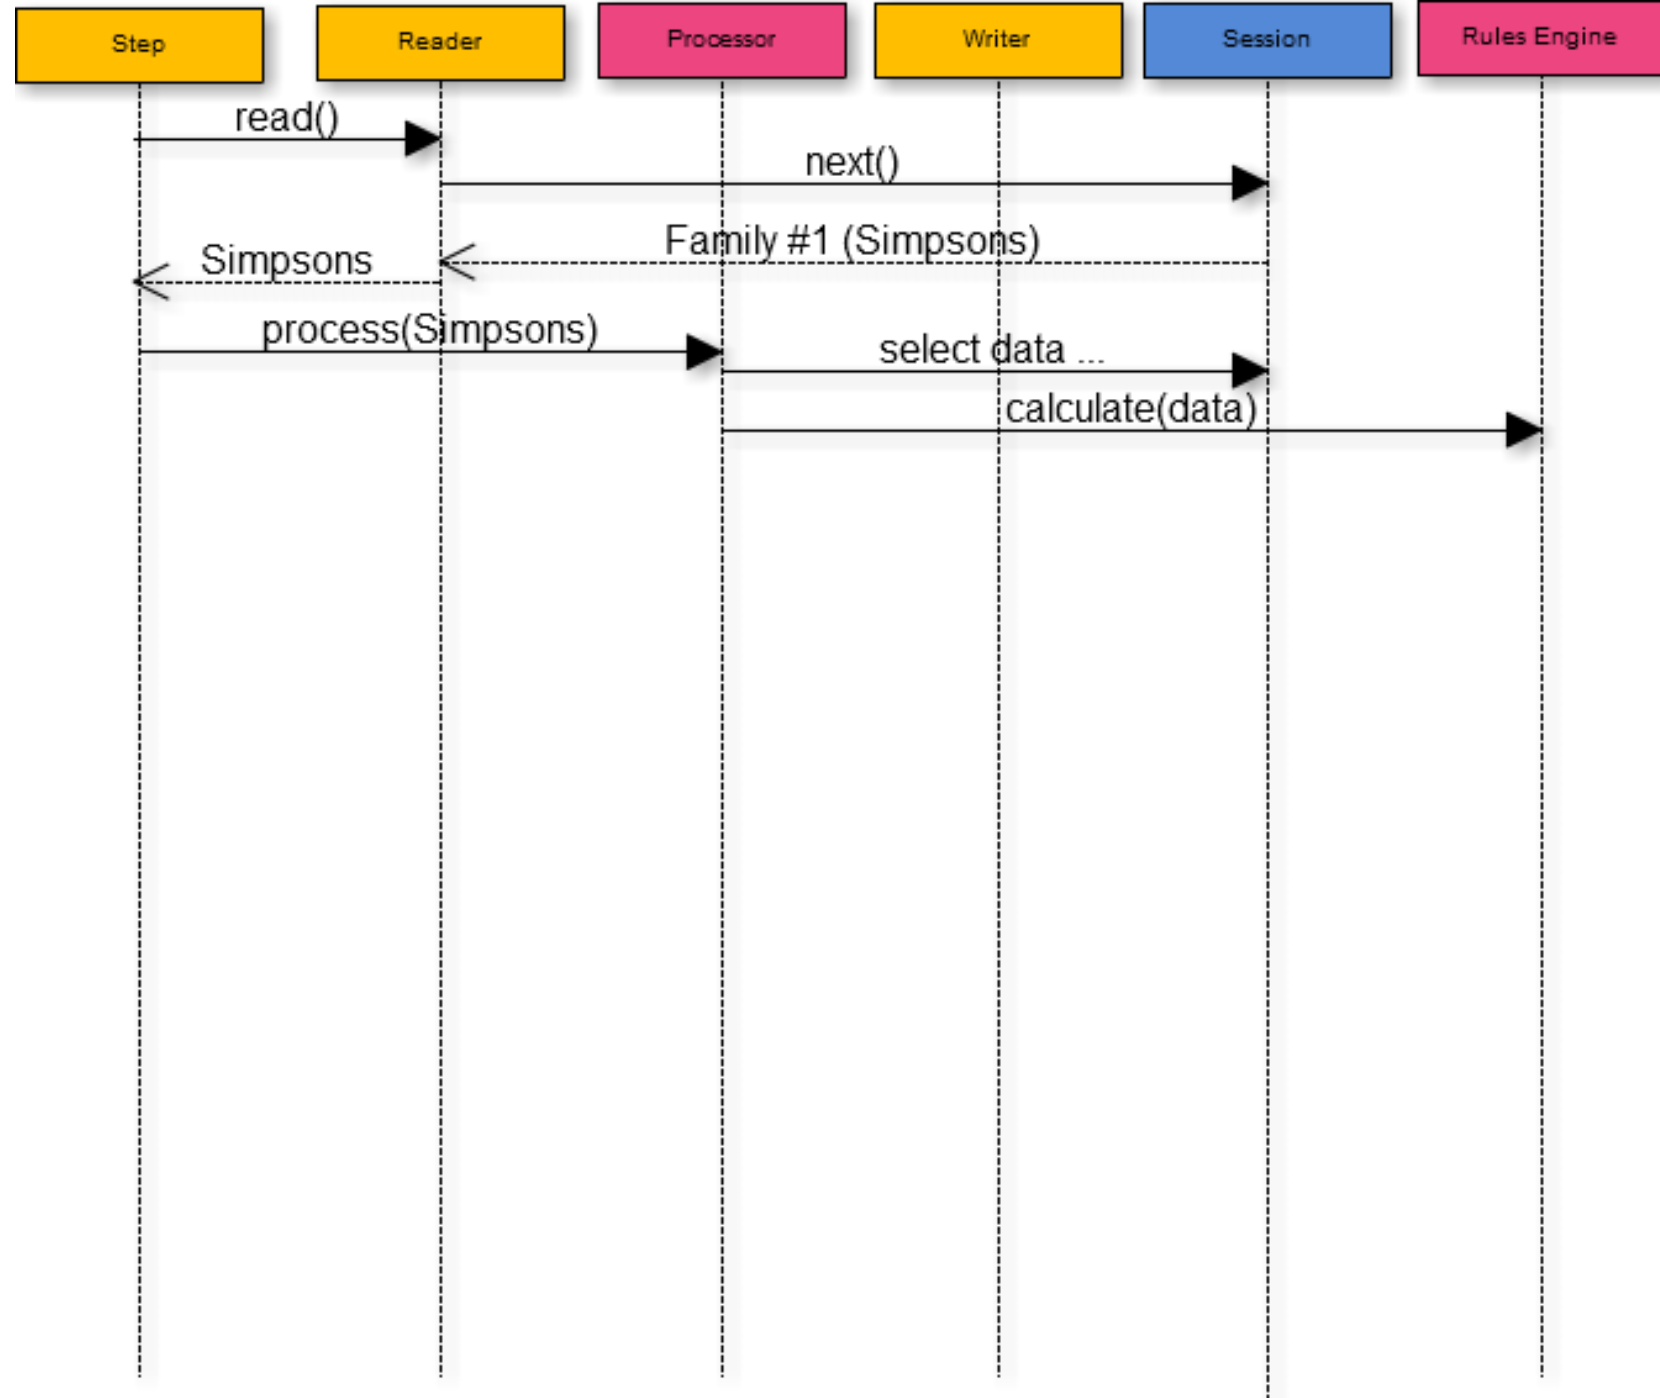


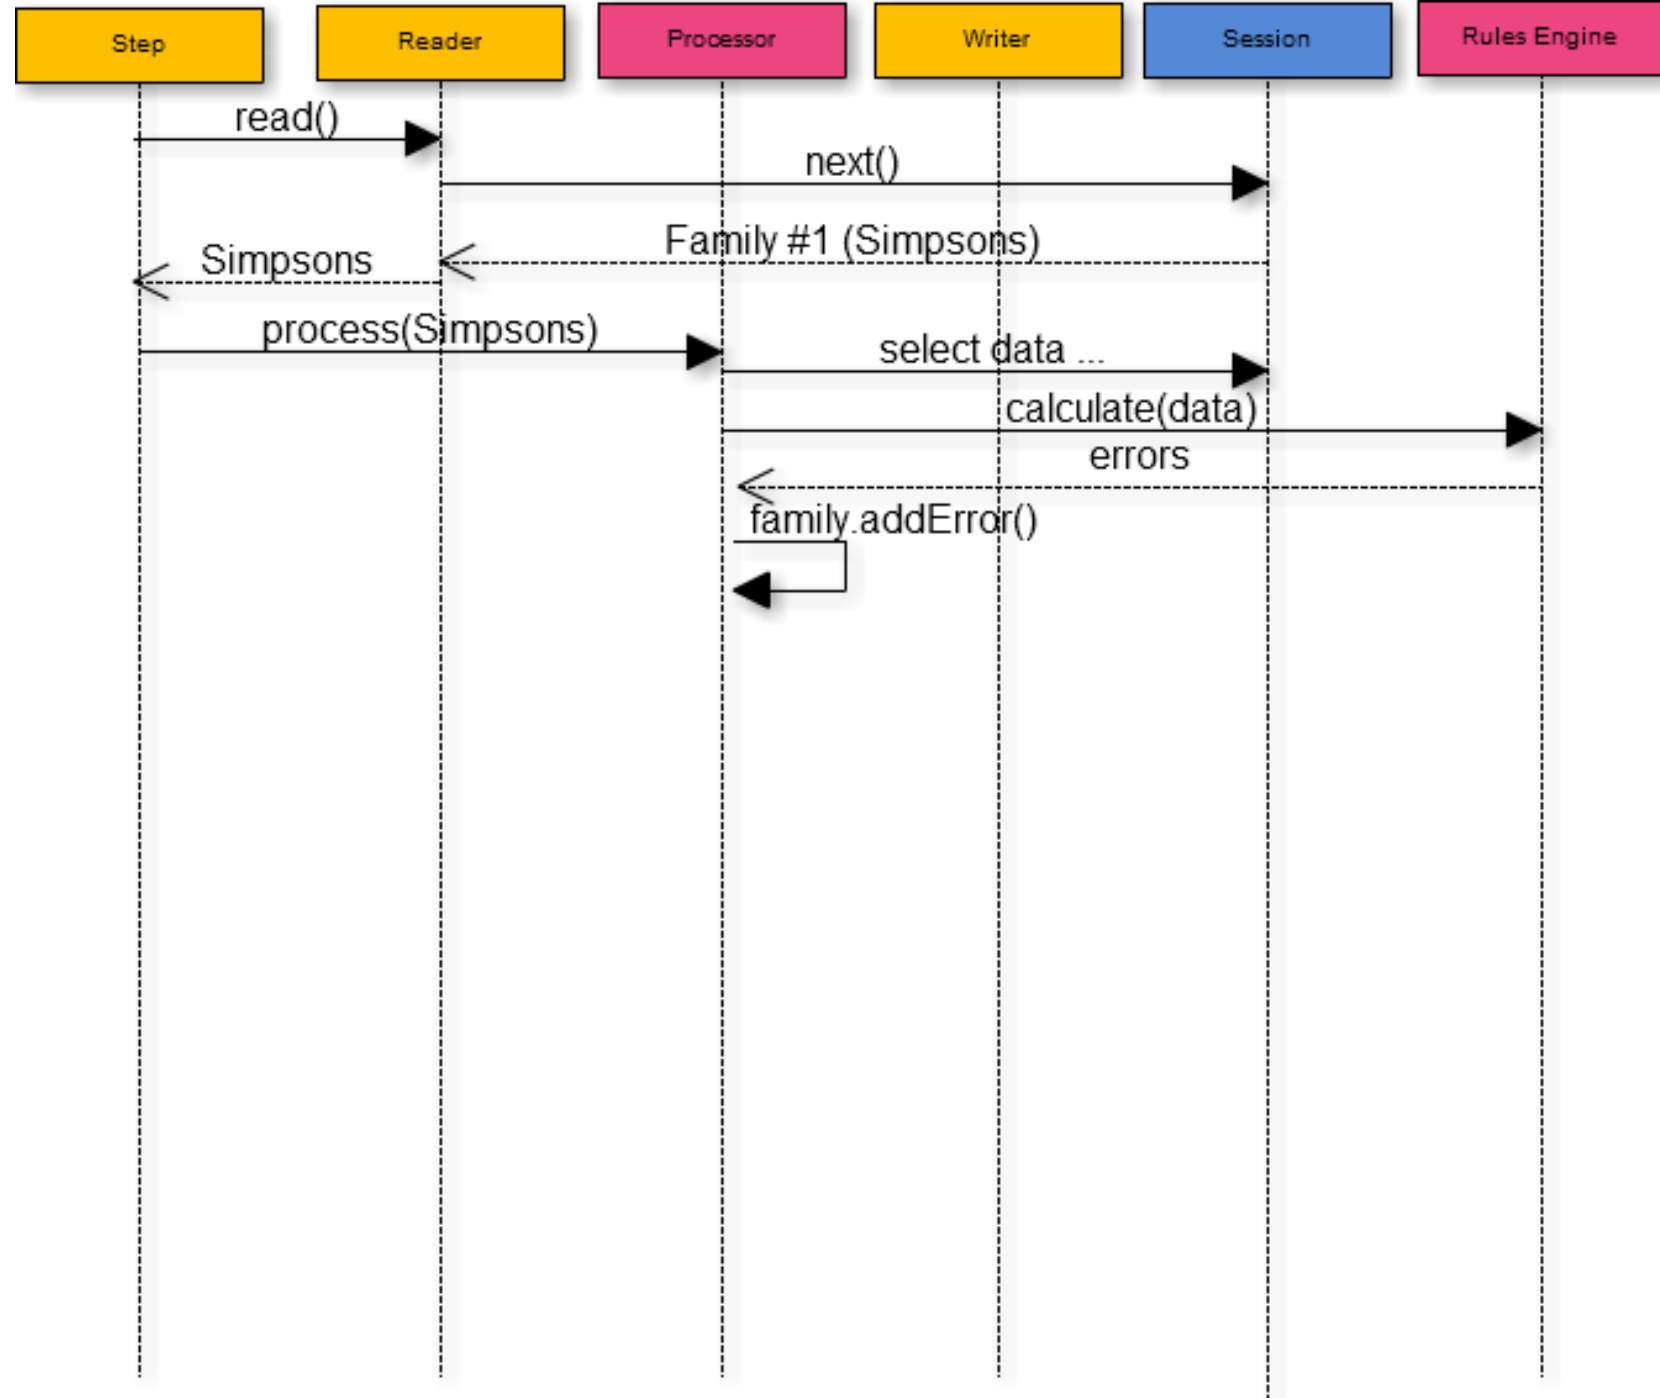


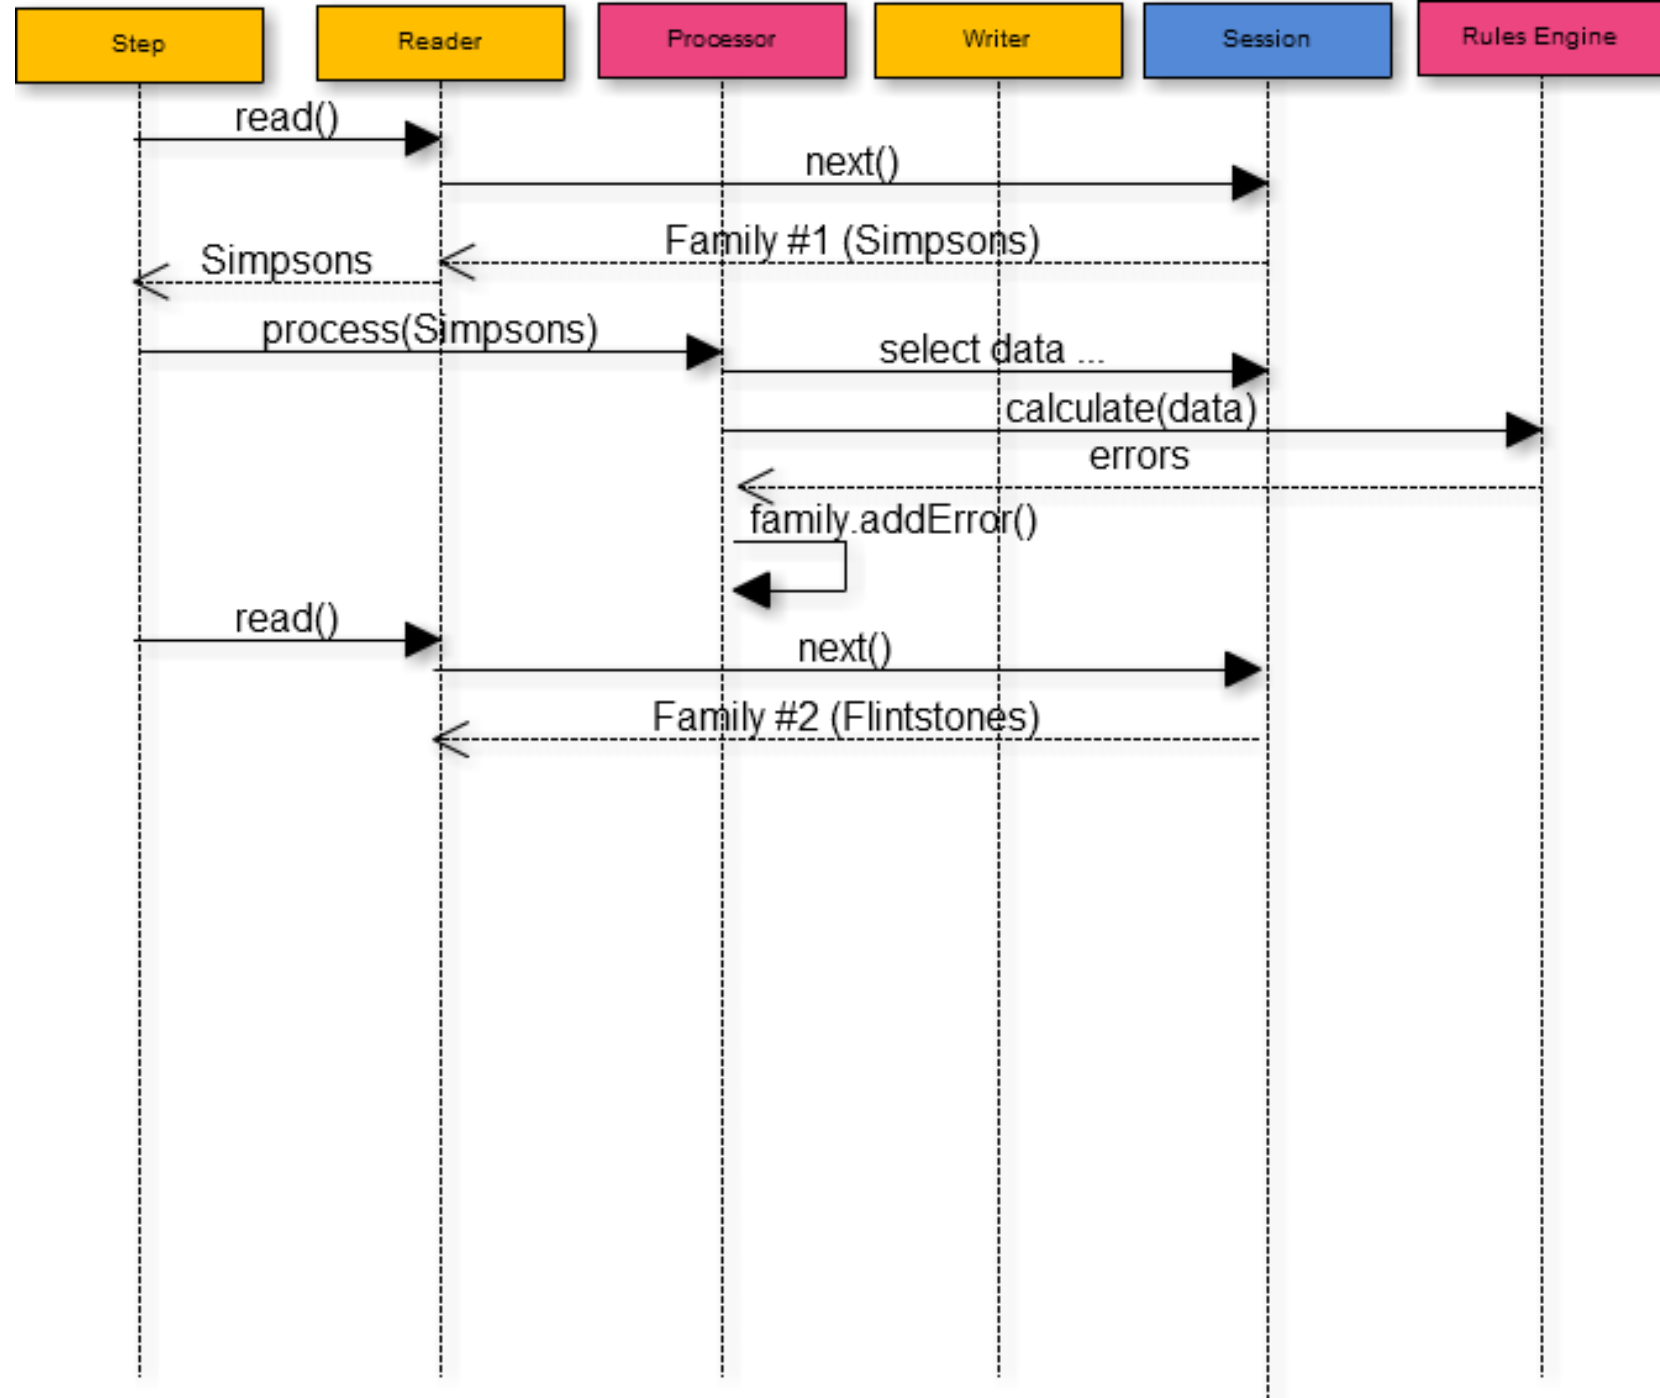


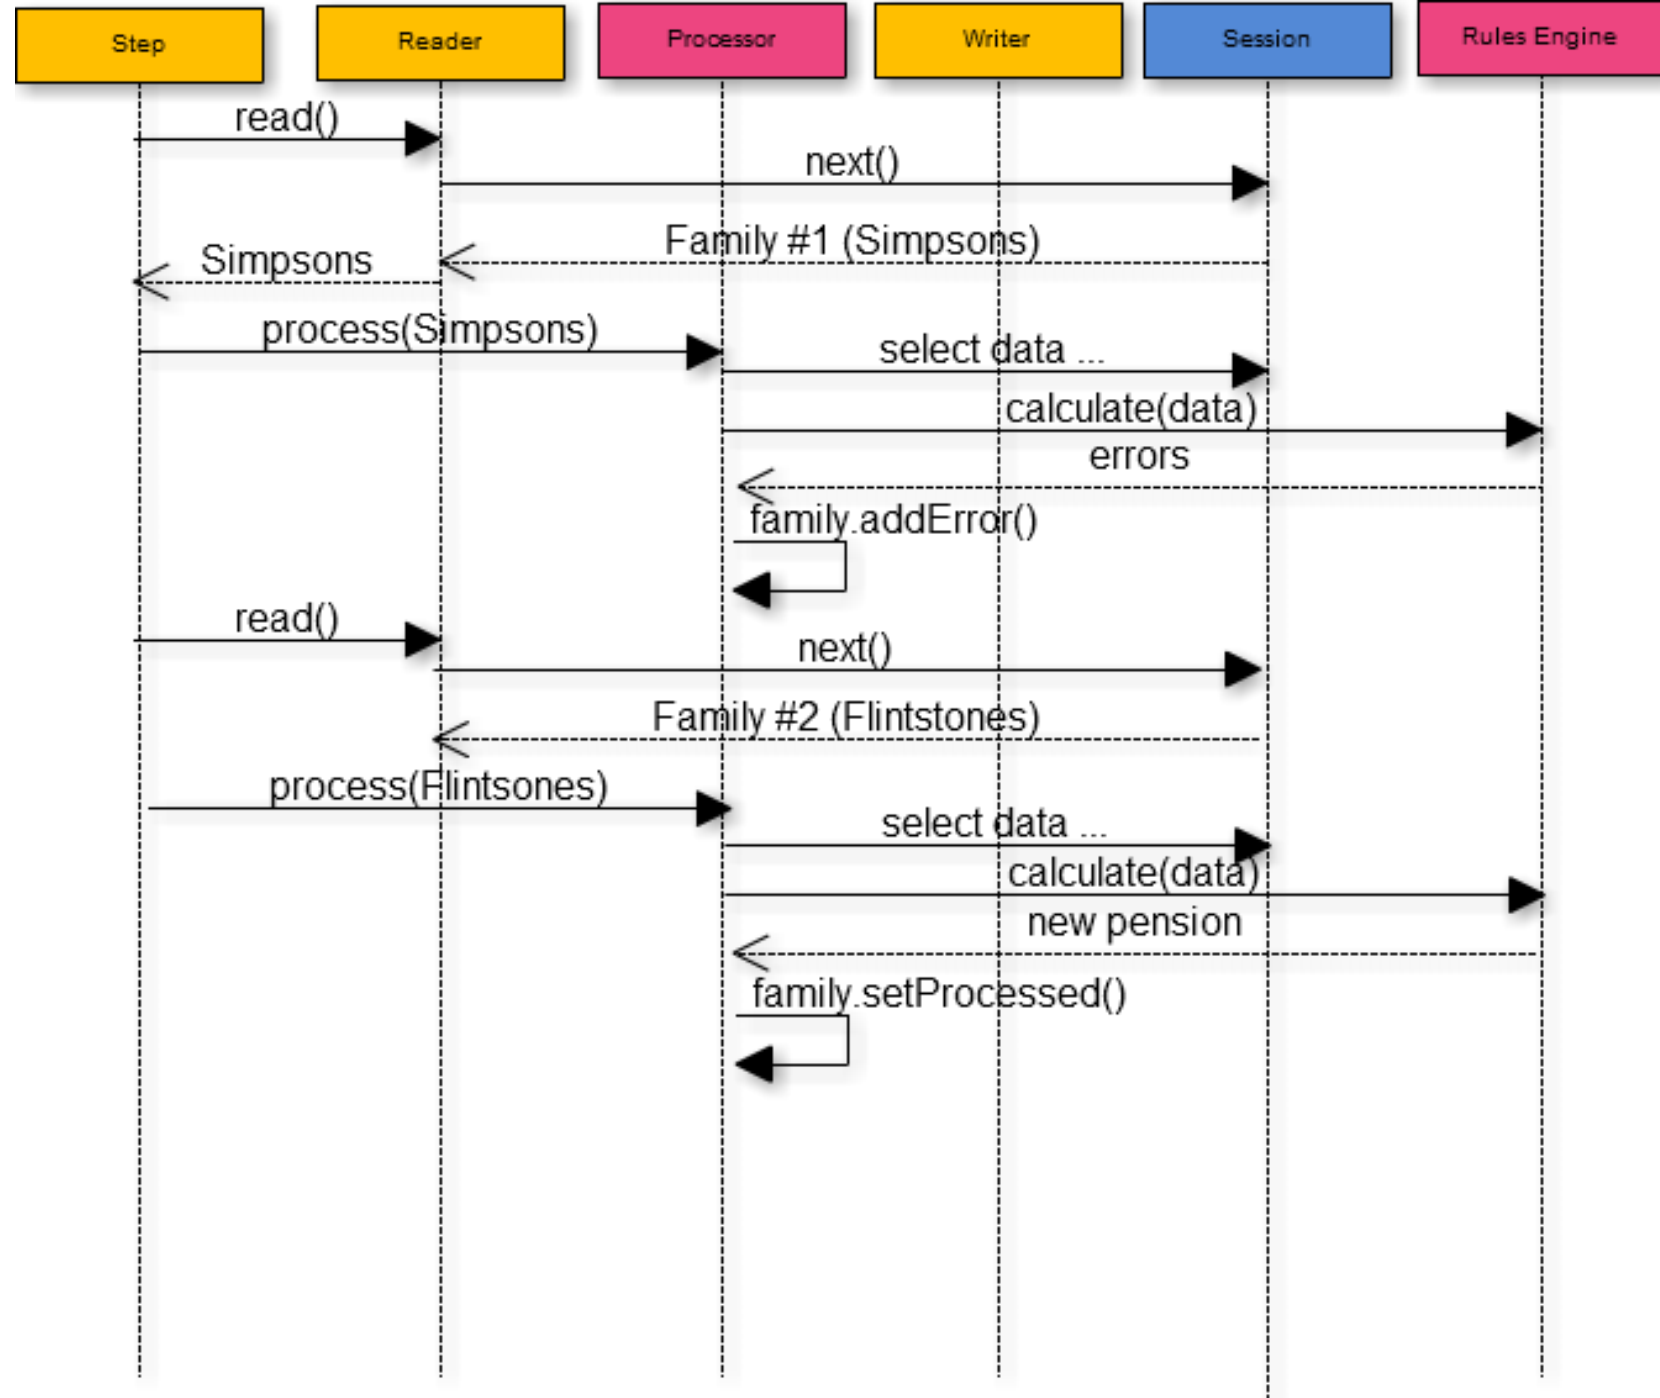


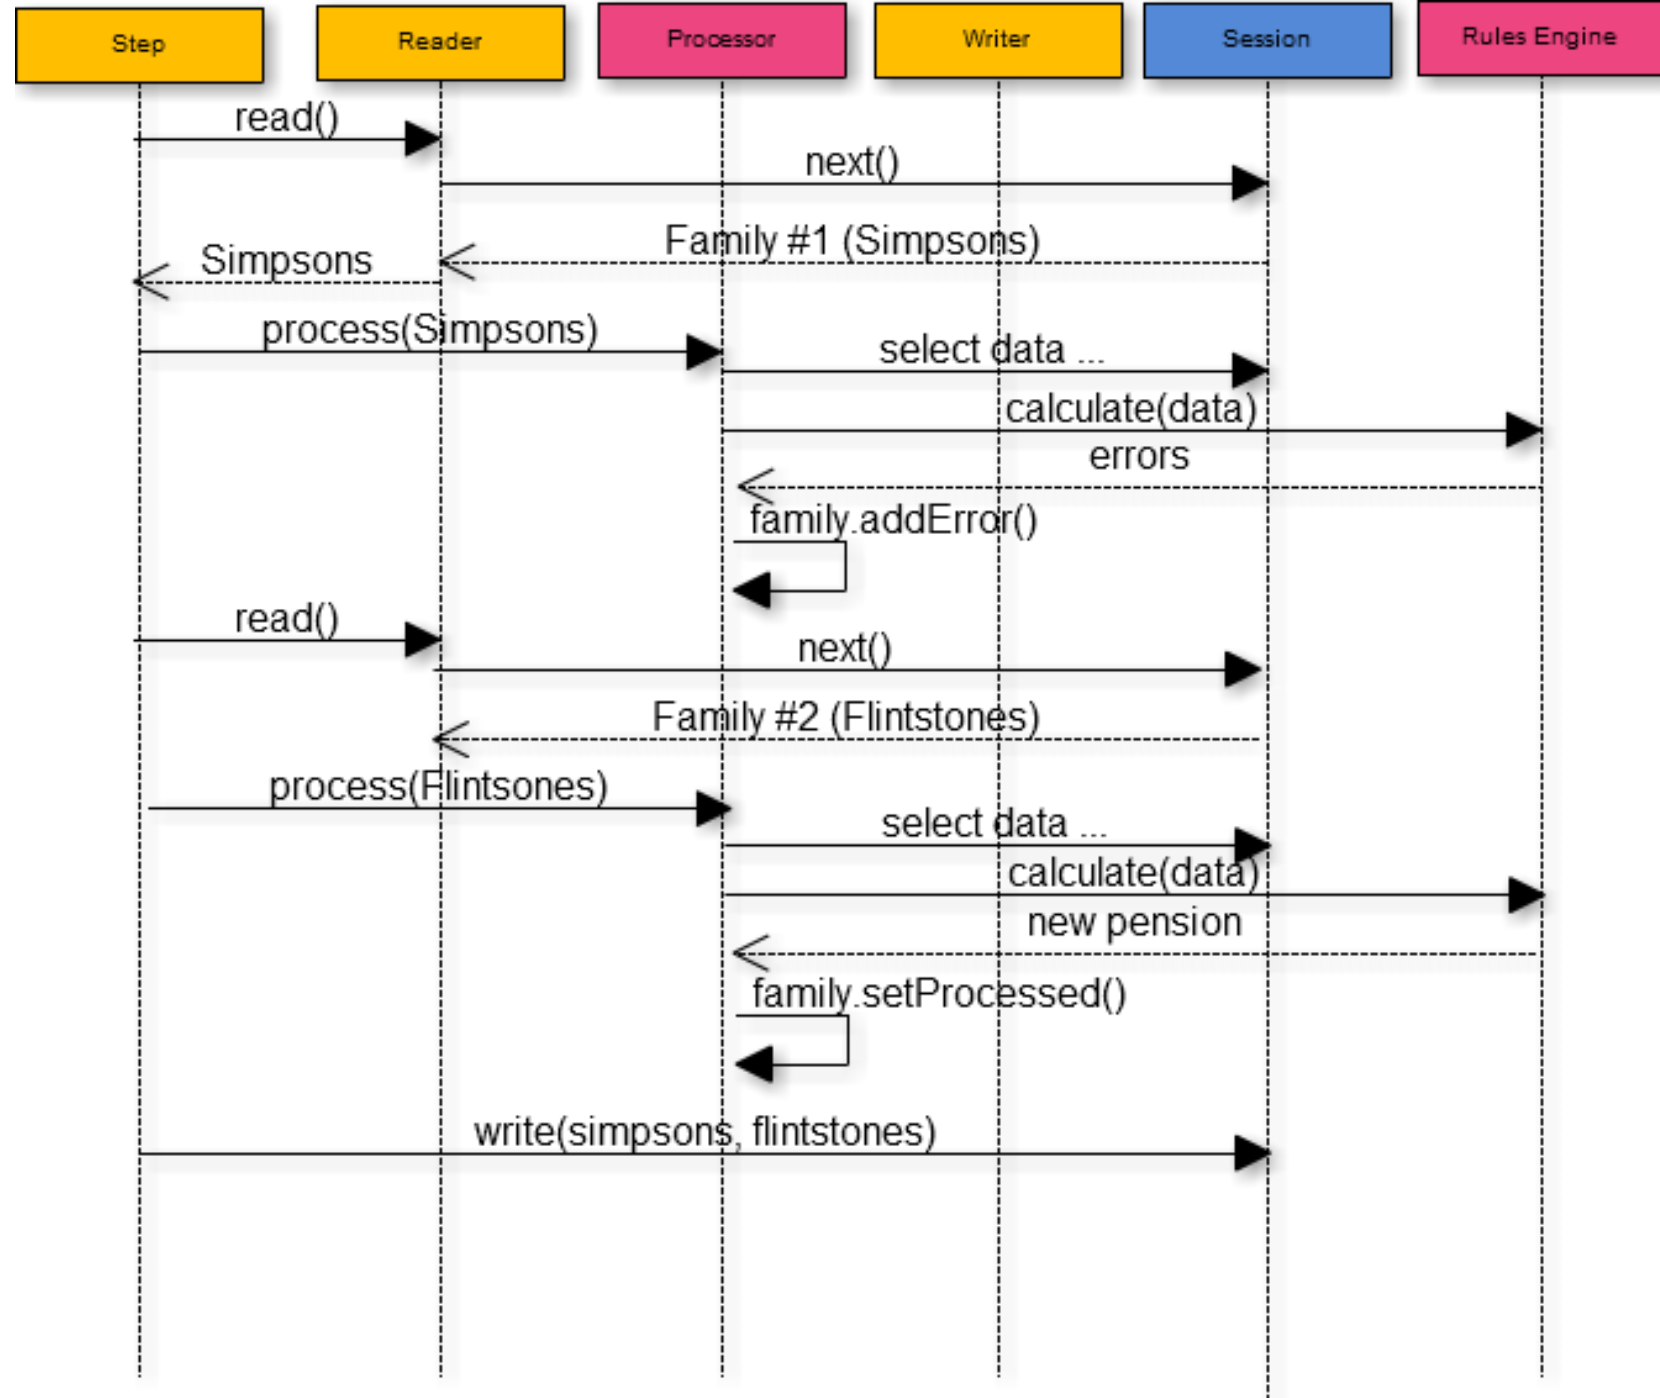




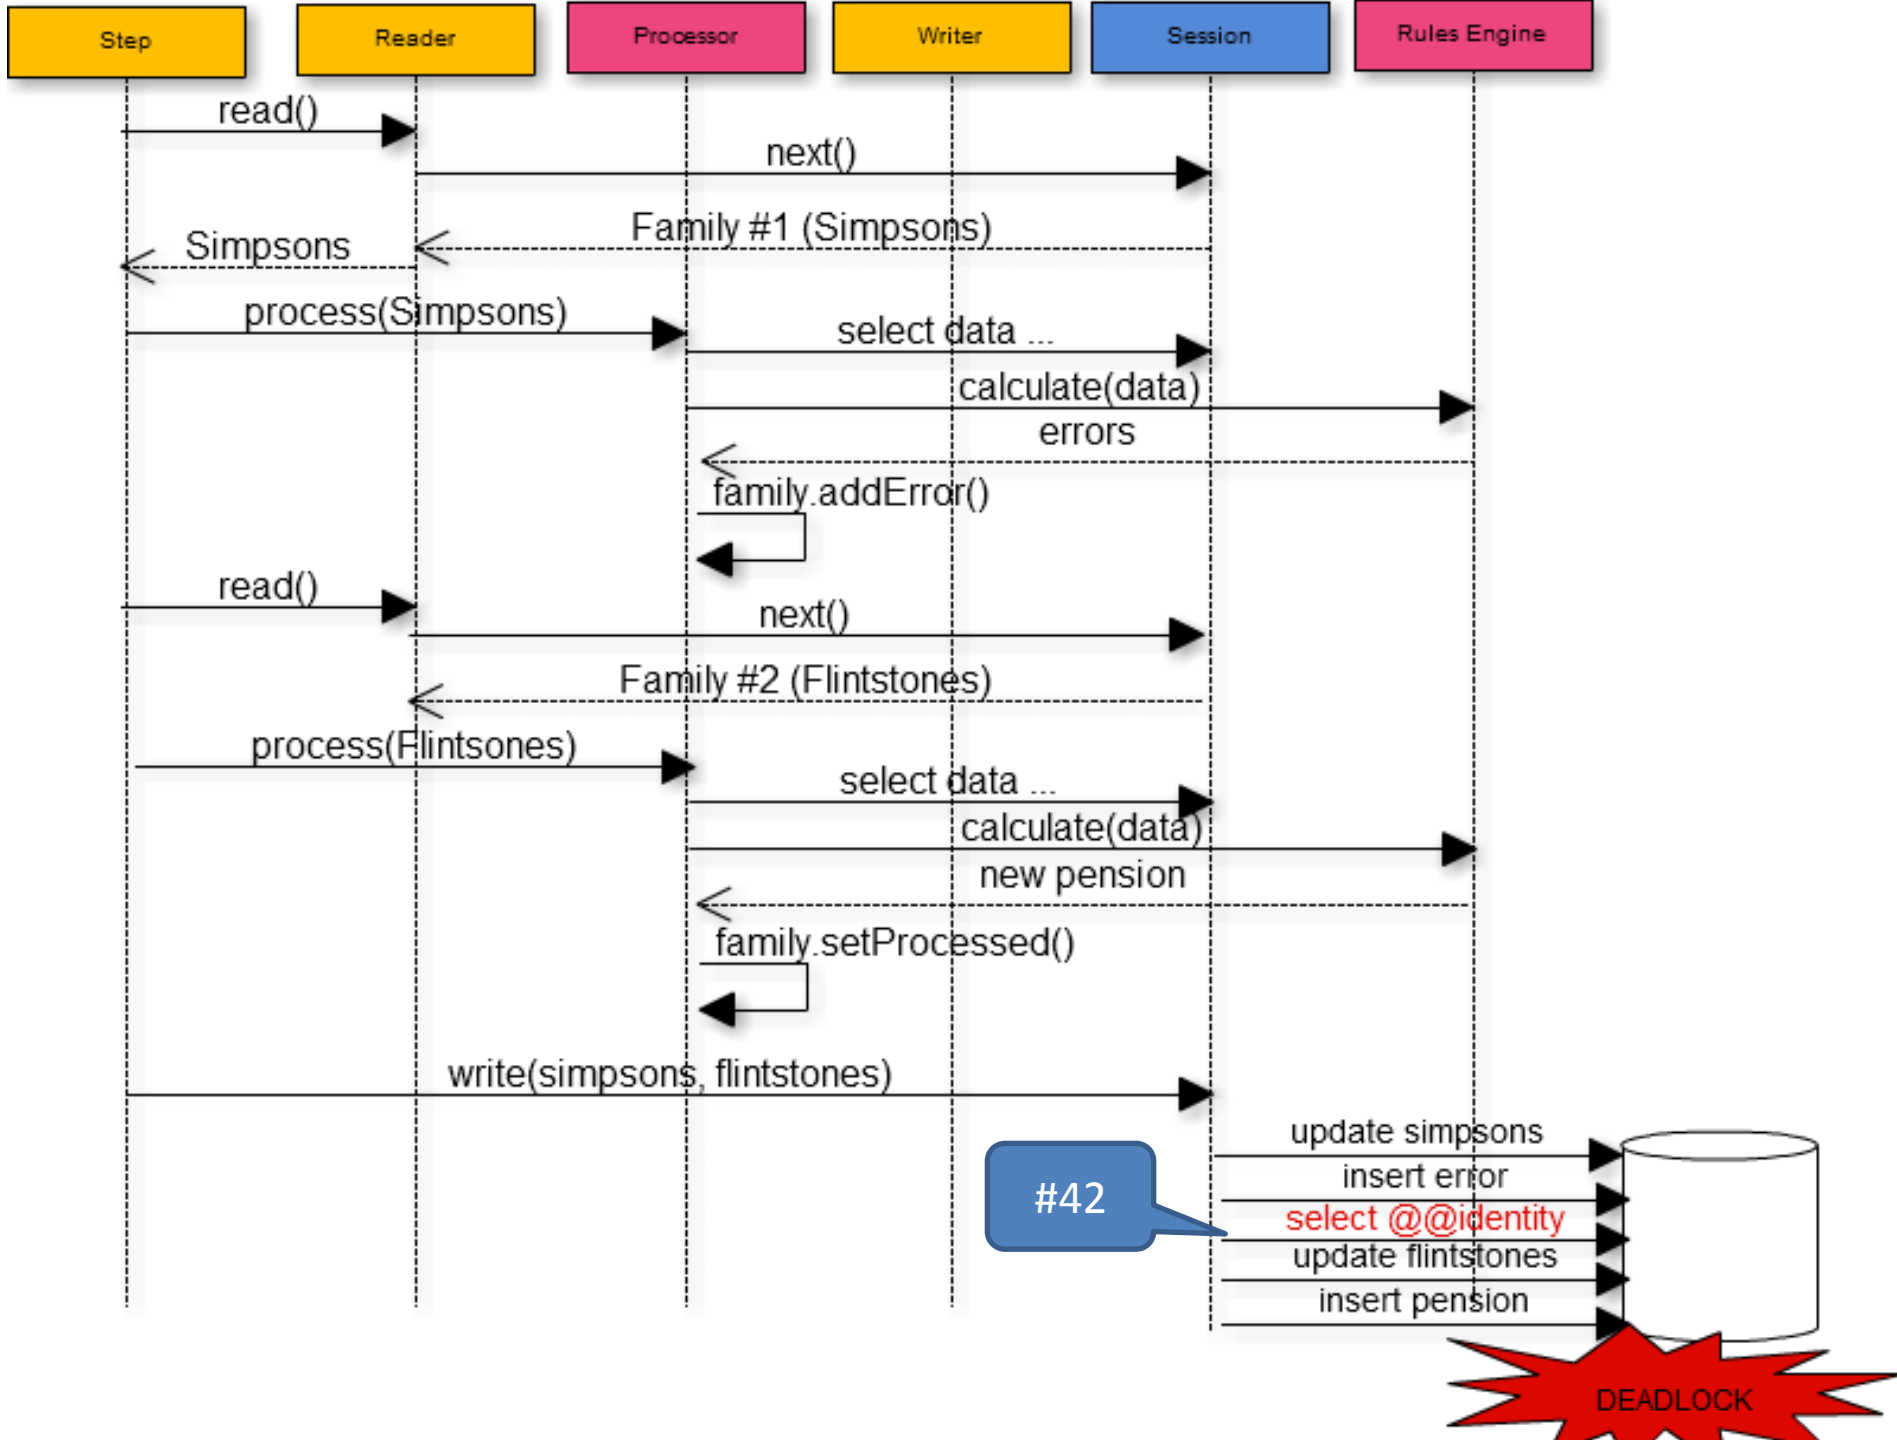








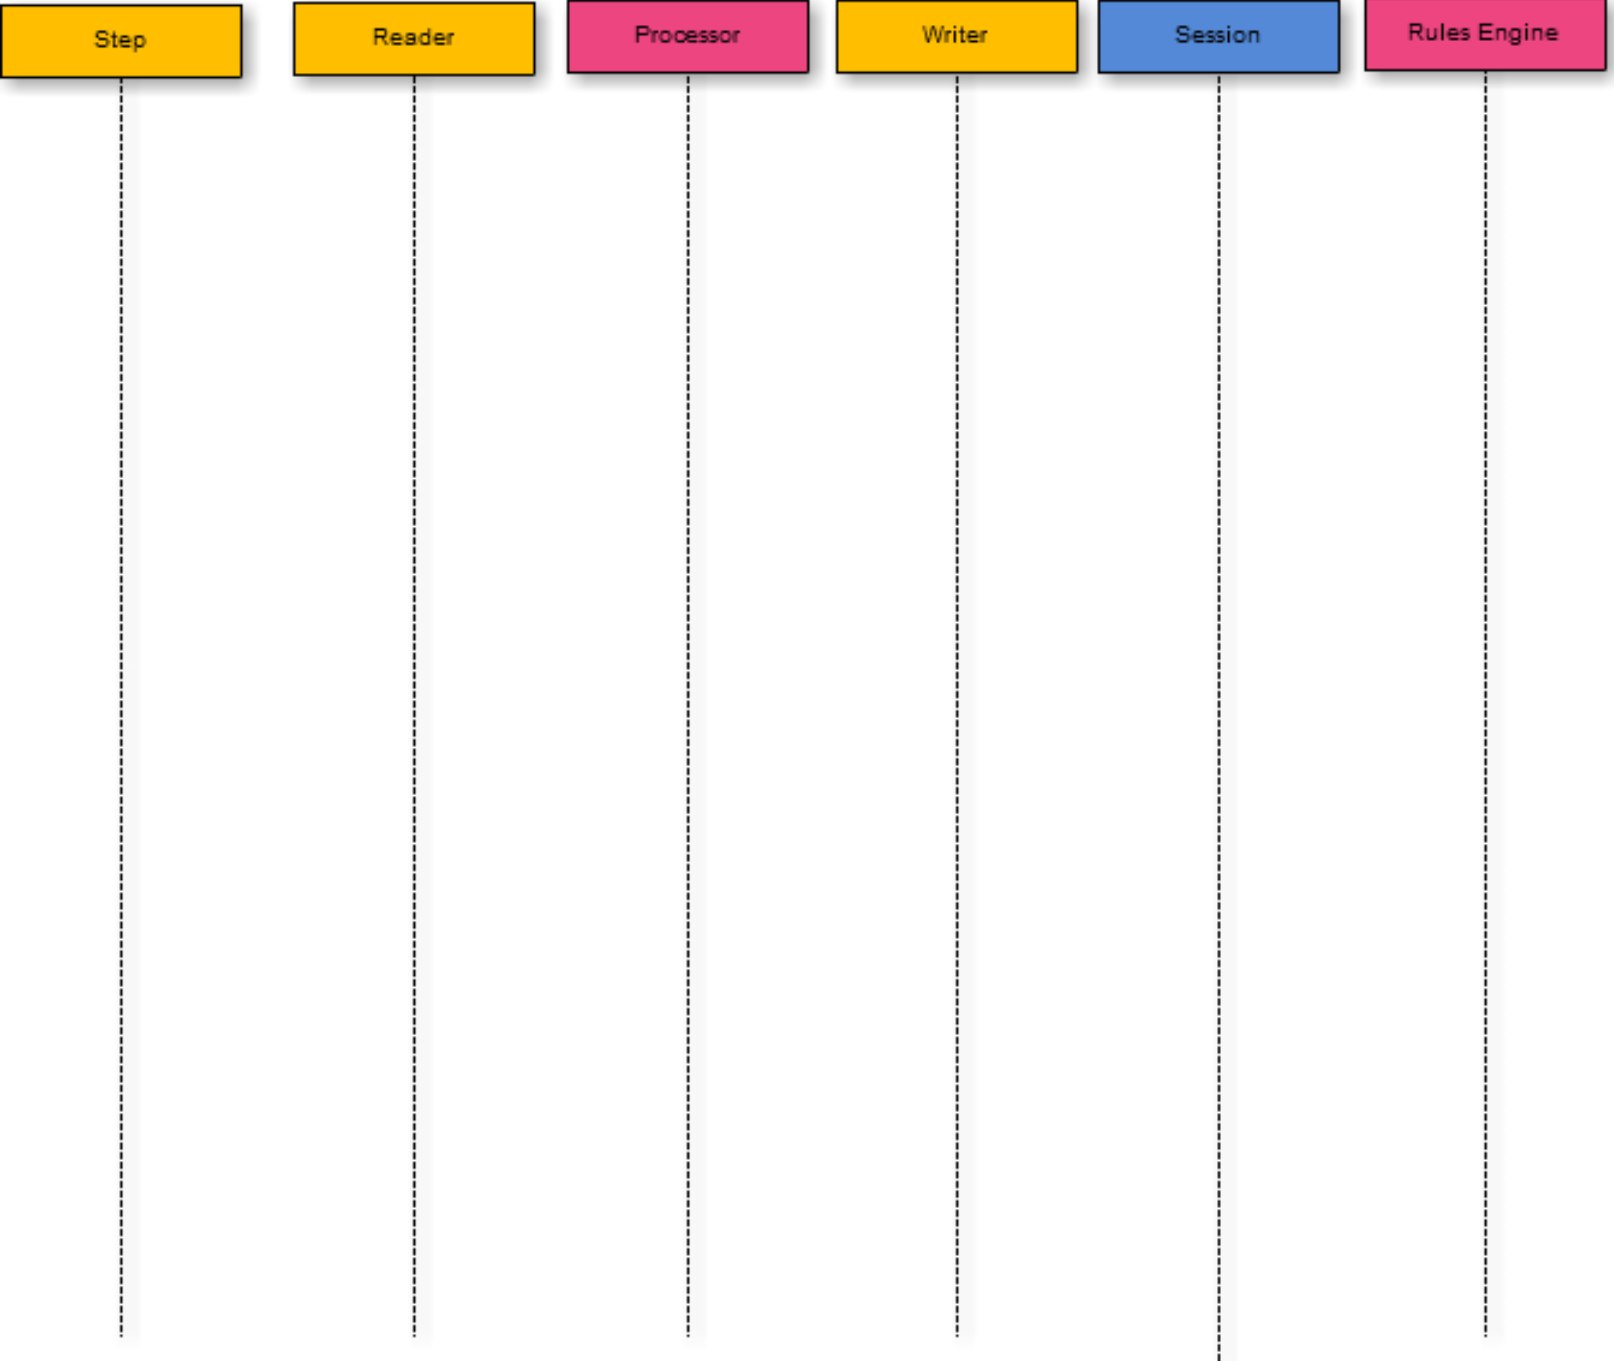


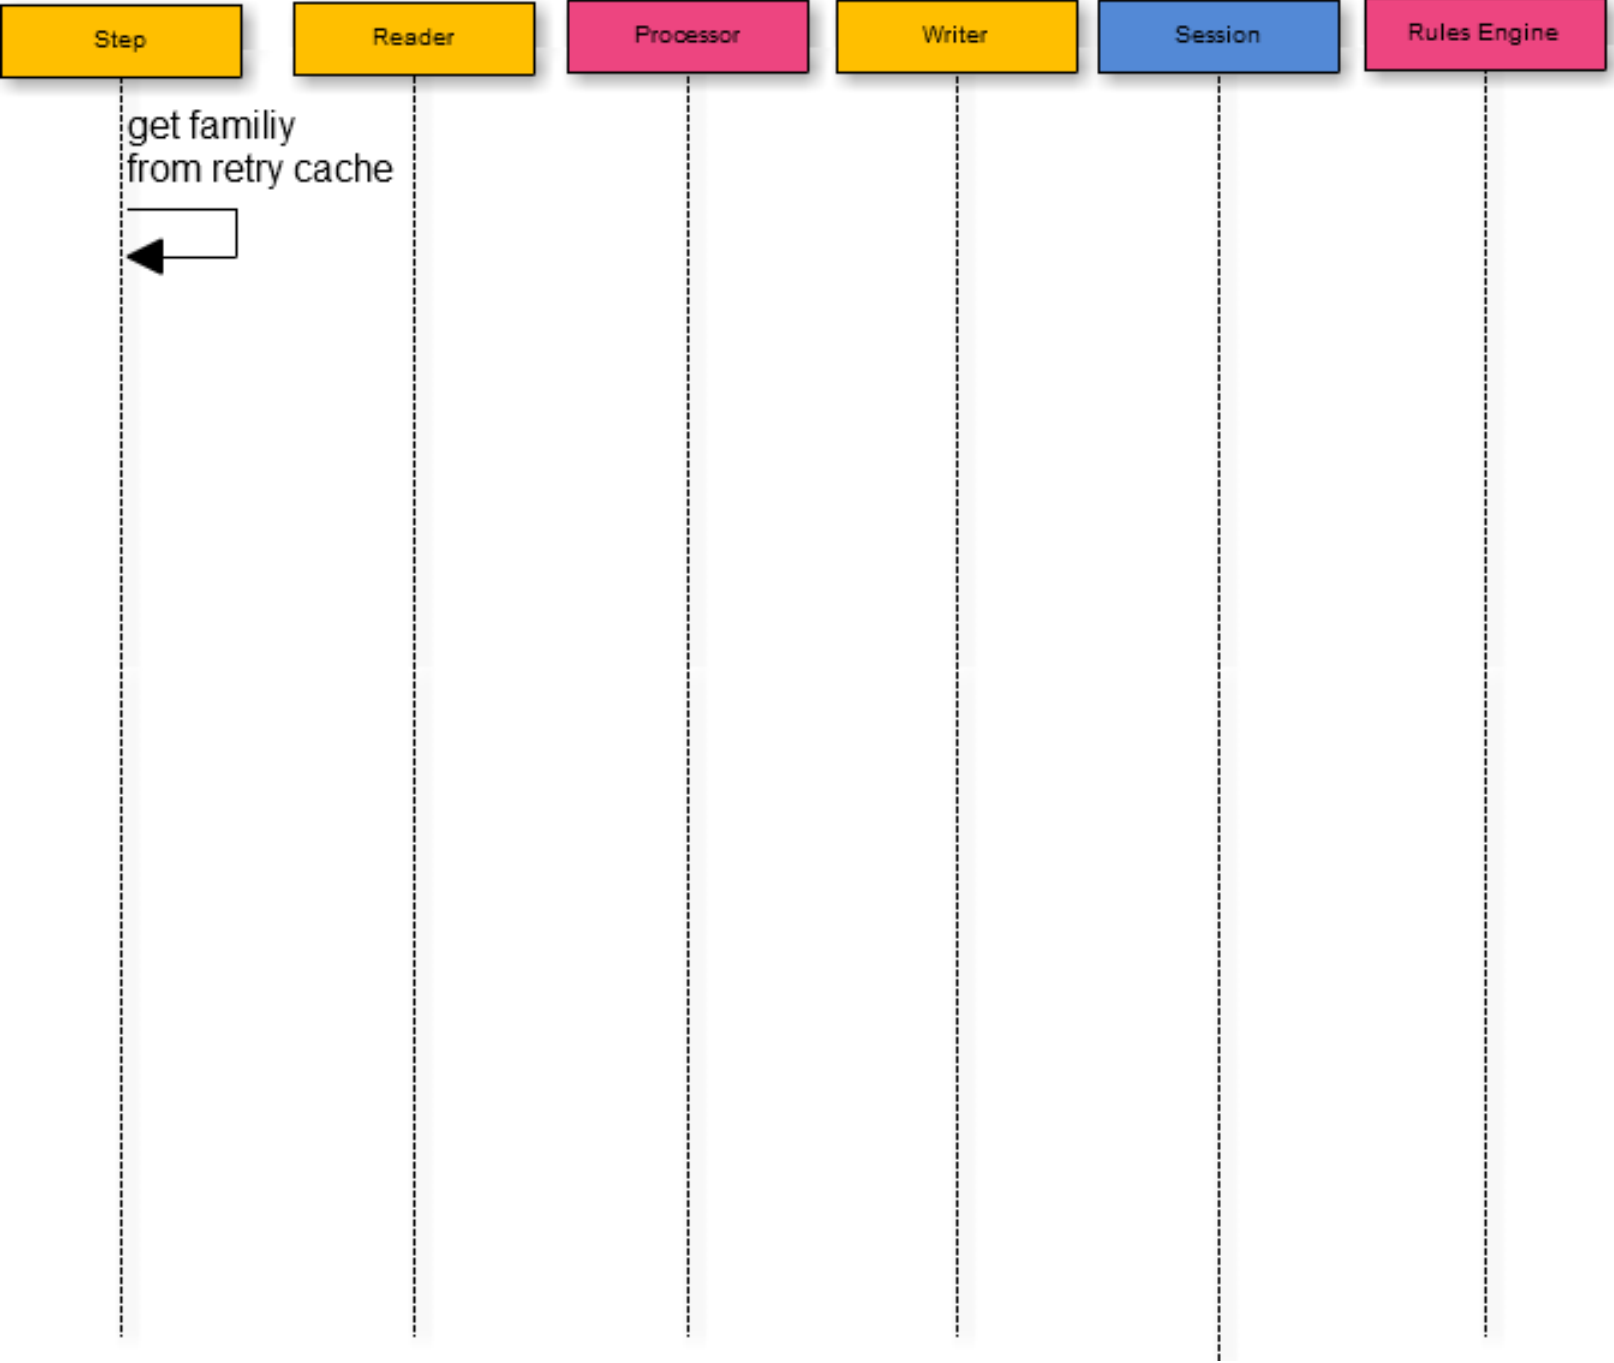


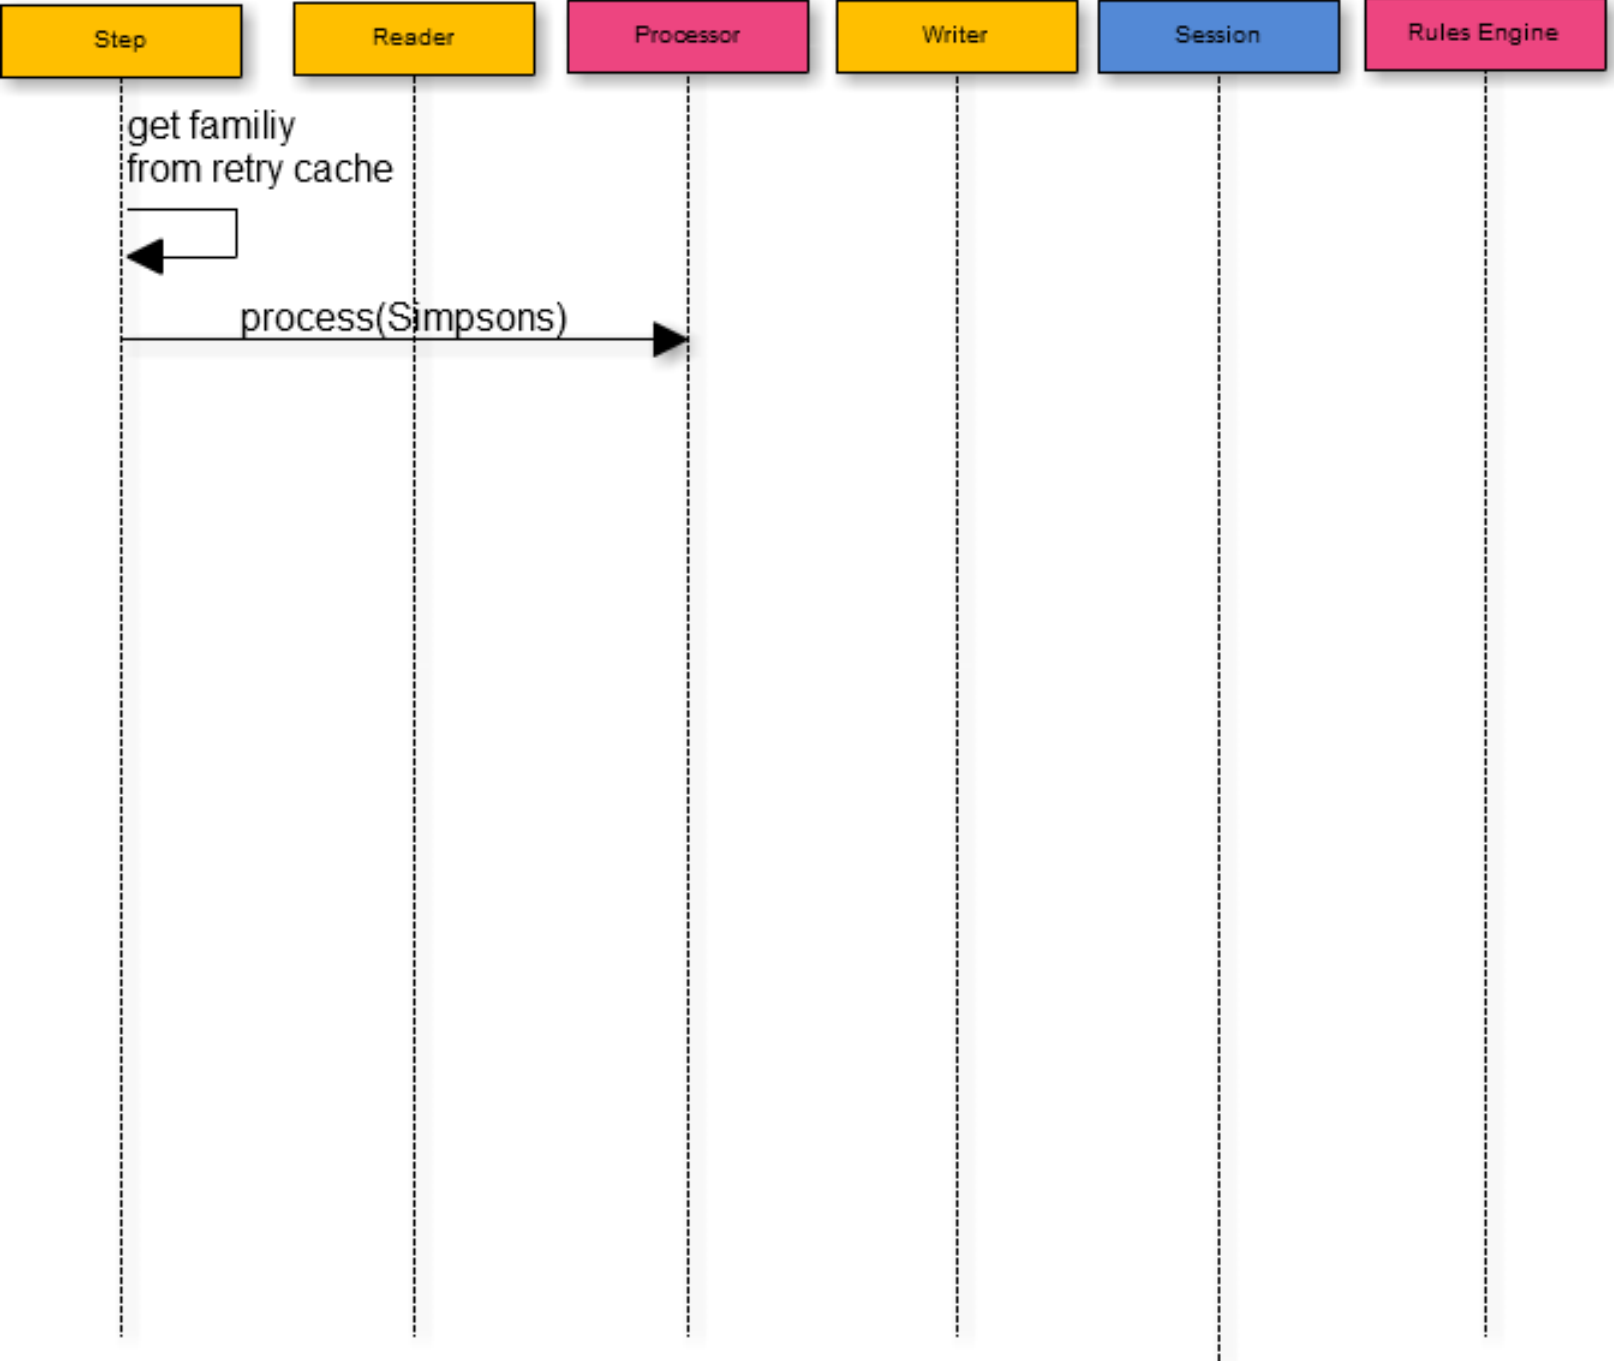


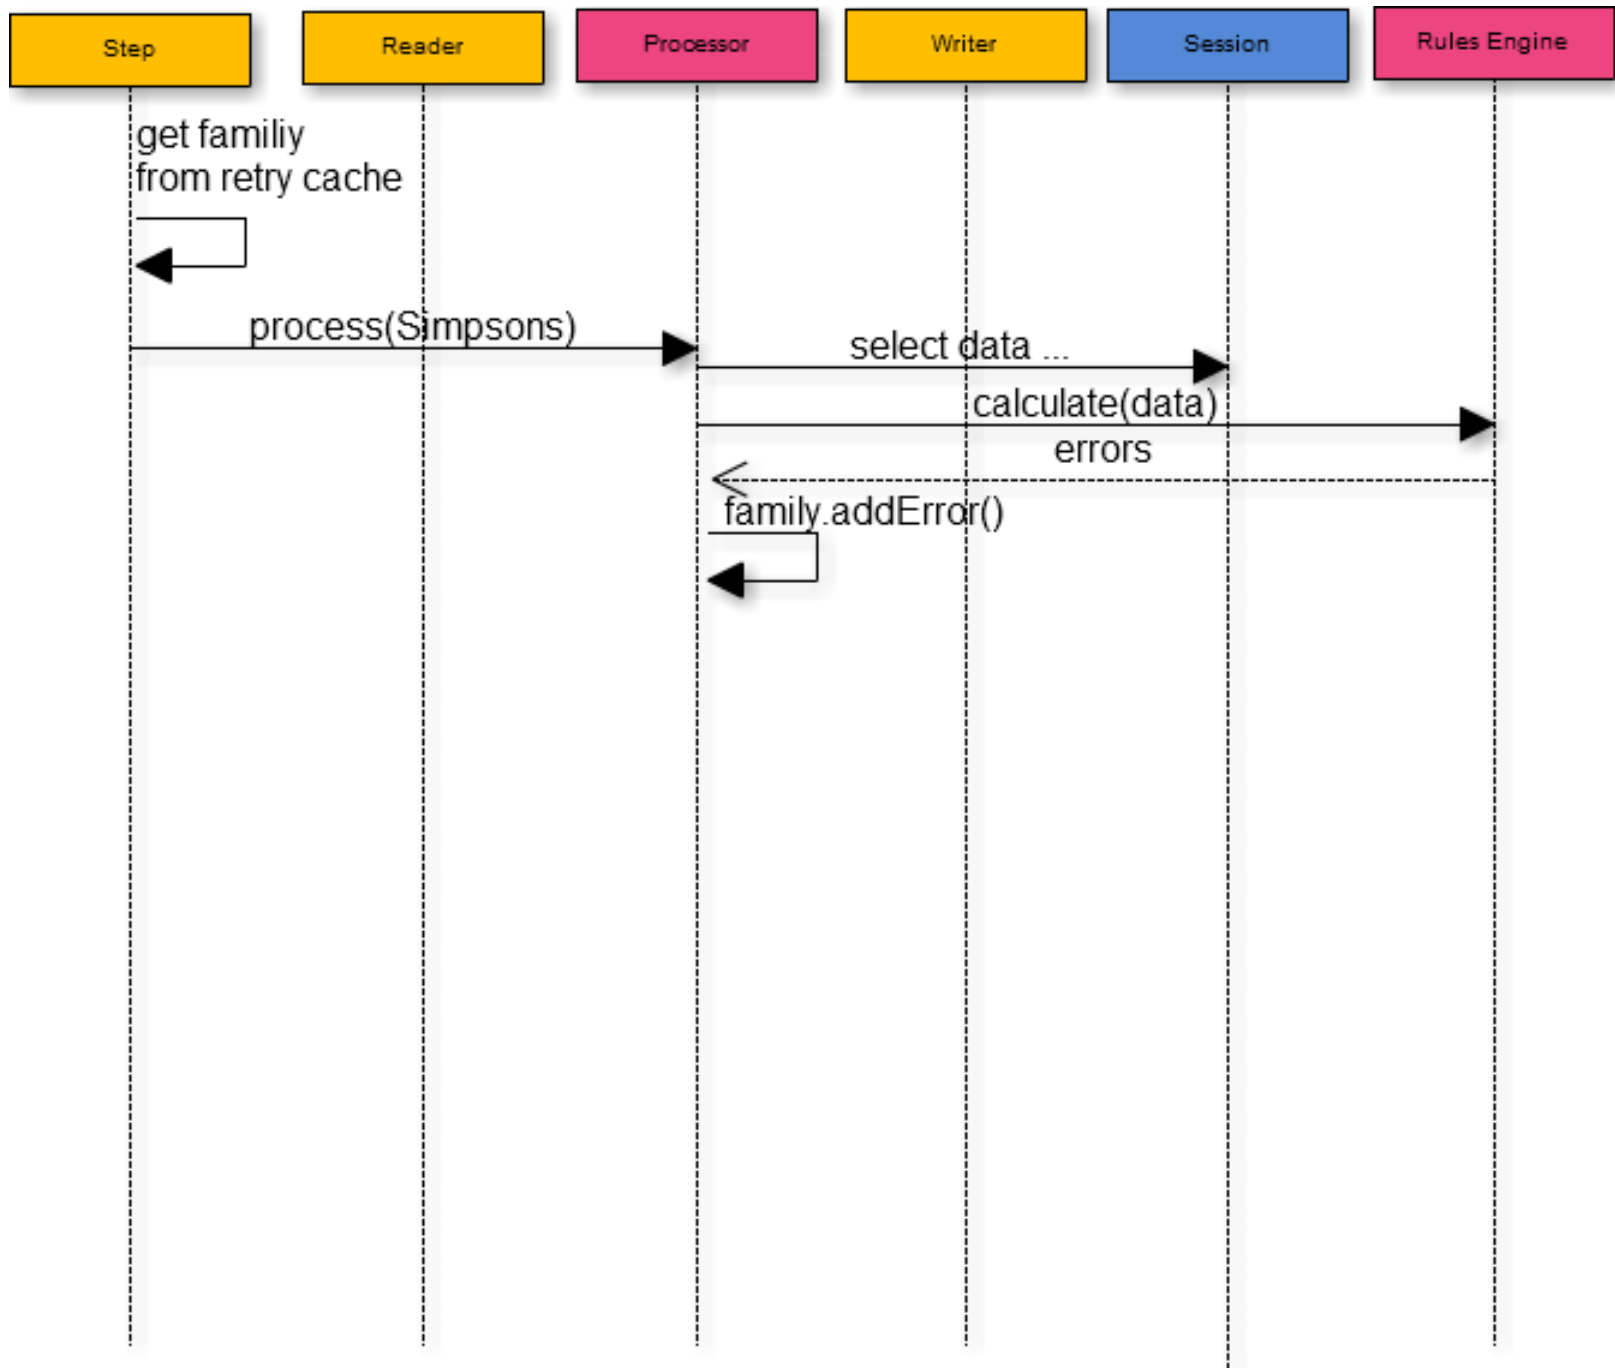
*High performance. Delivered.*

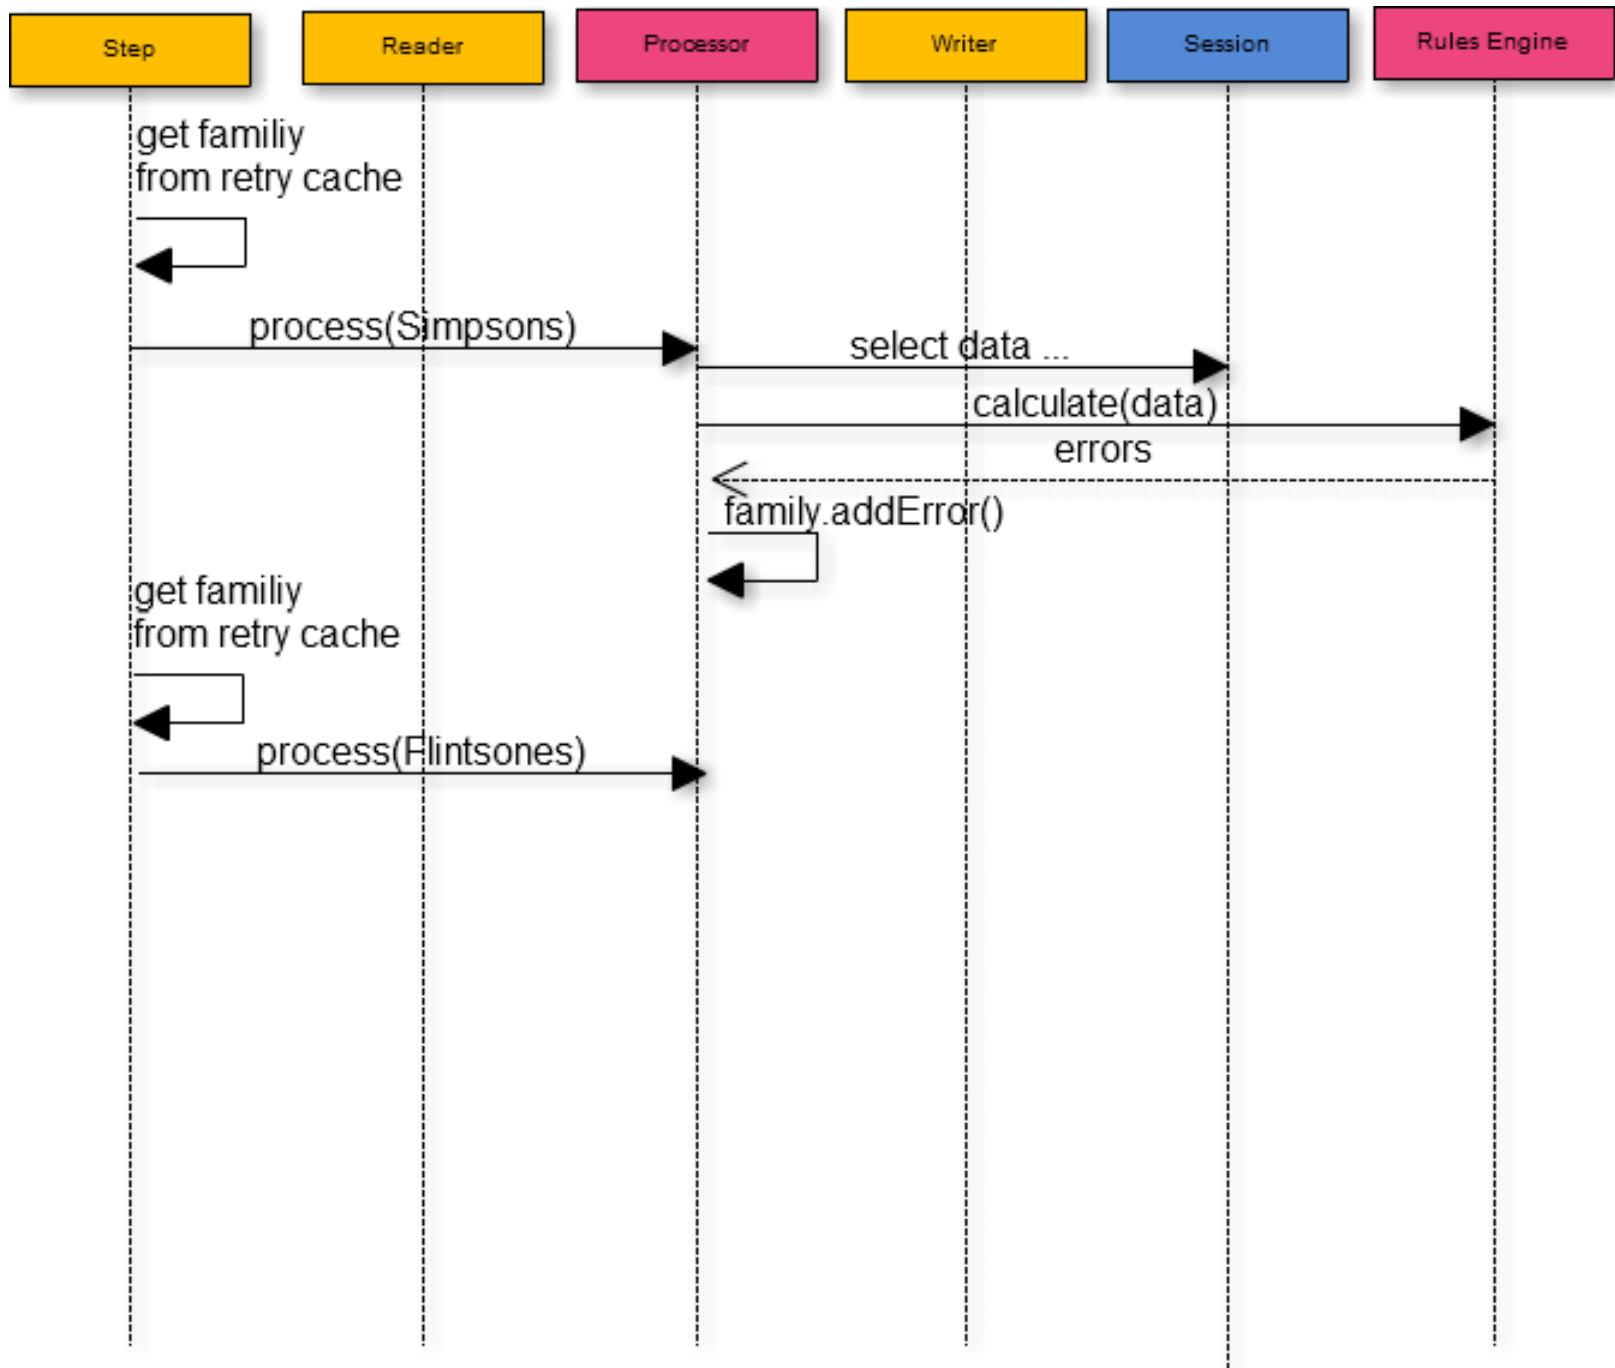
# Retry

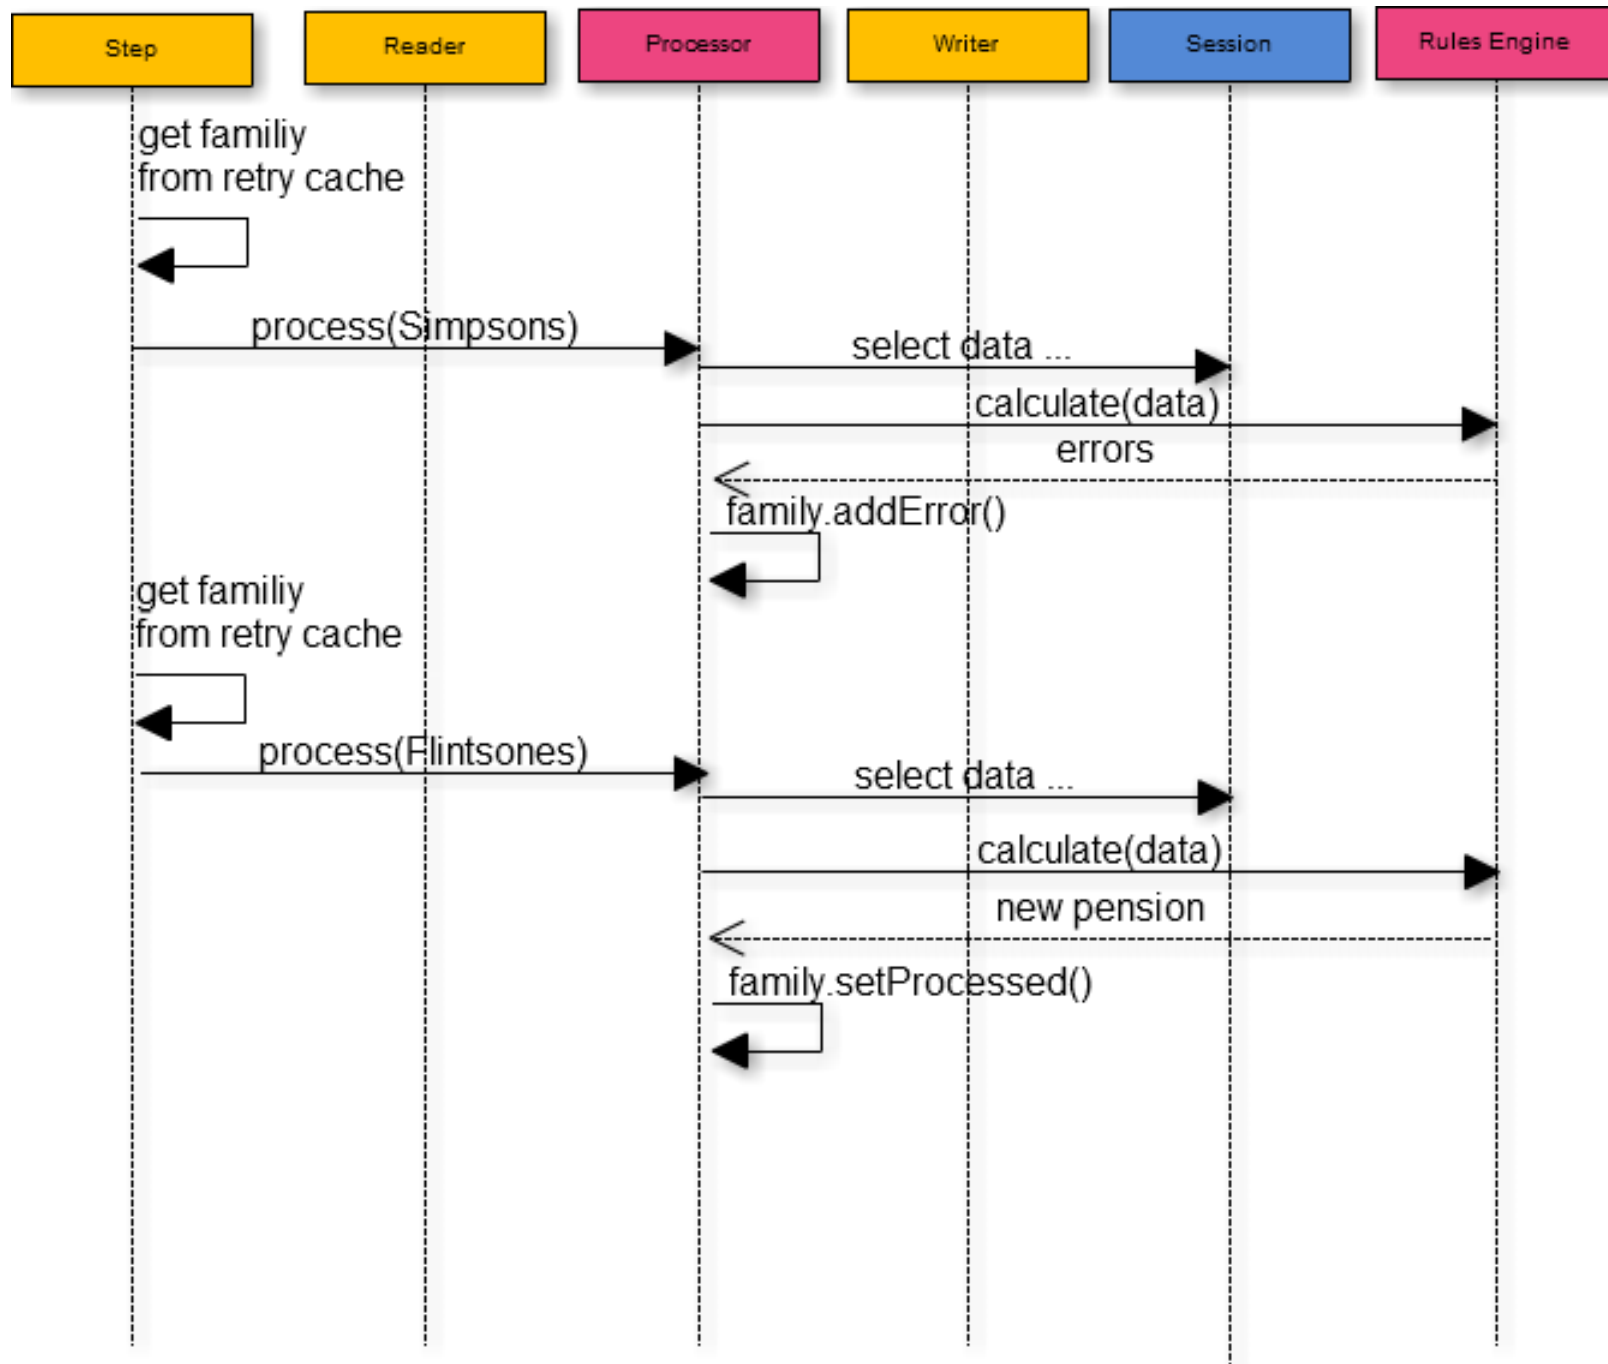




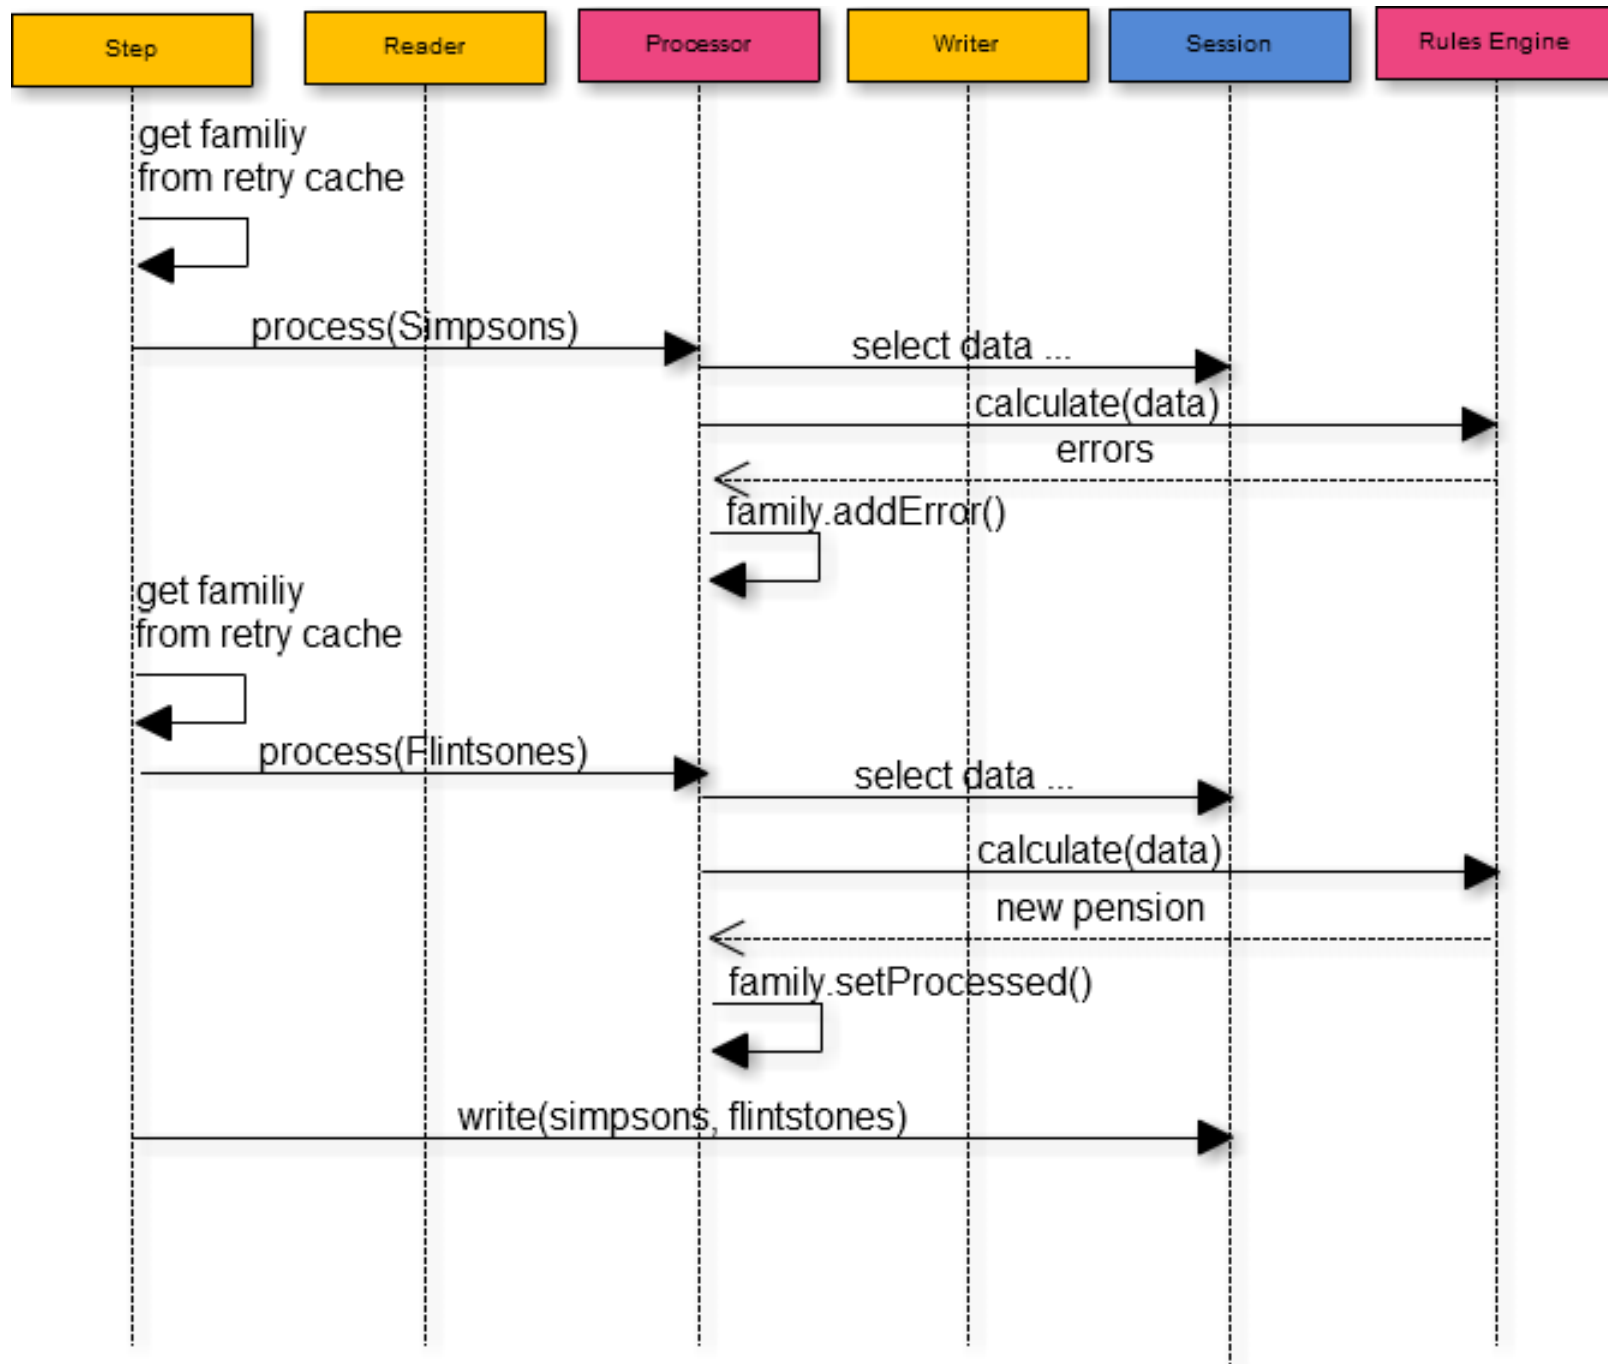


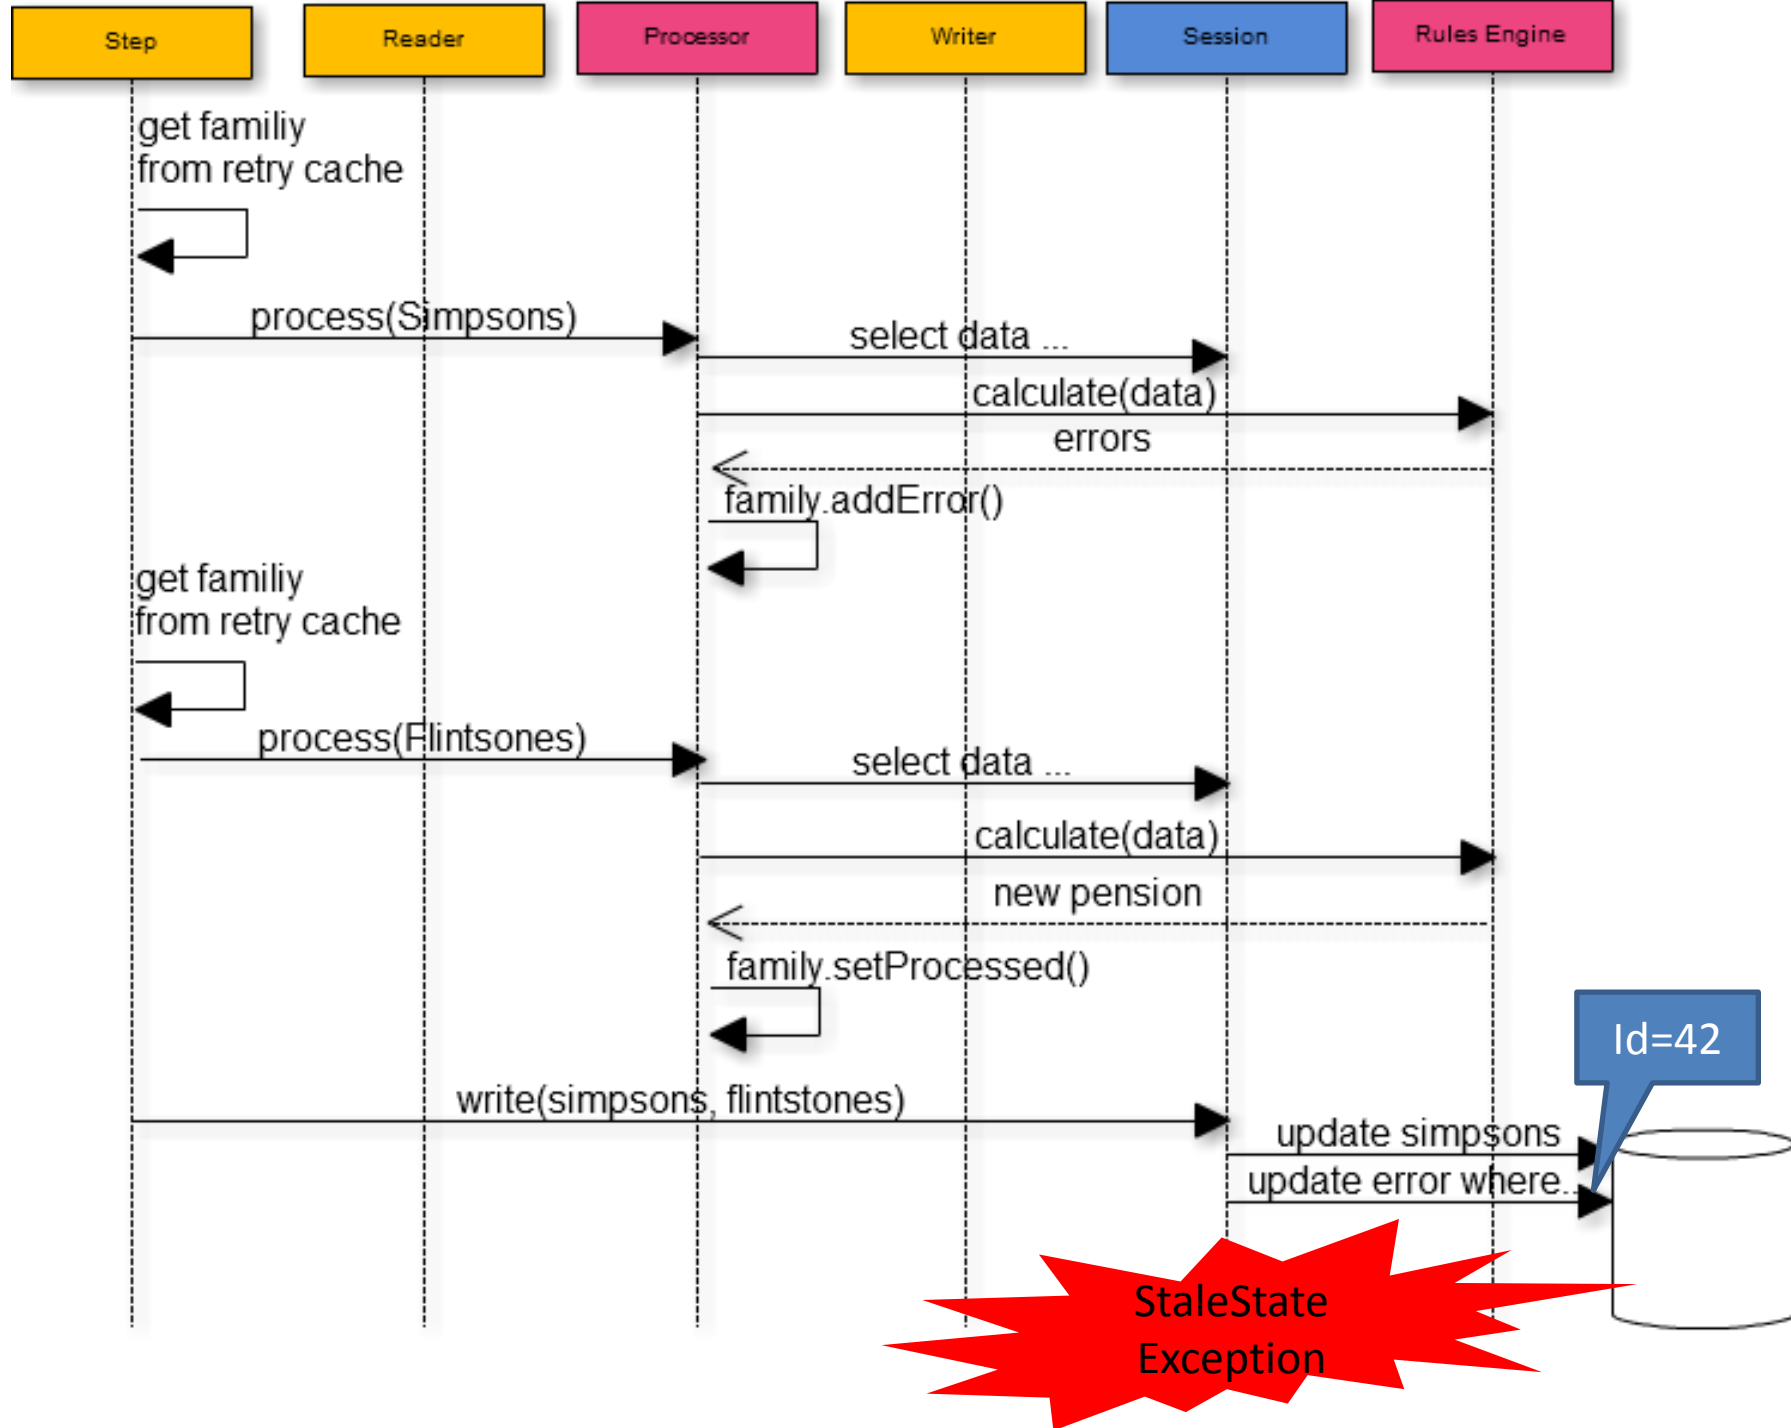














*High performance. Delivered.*

# What do we do?



*High performance. Delivered.*

# **What we should have done weeks ago**

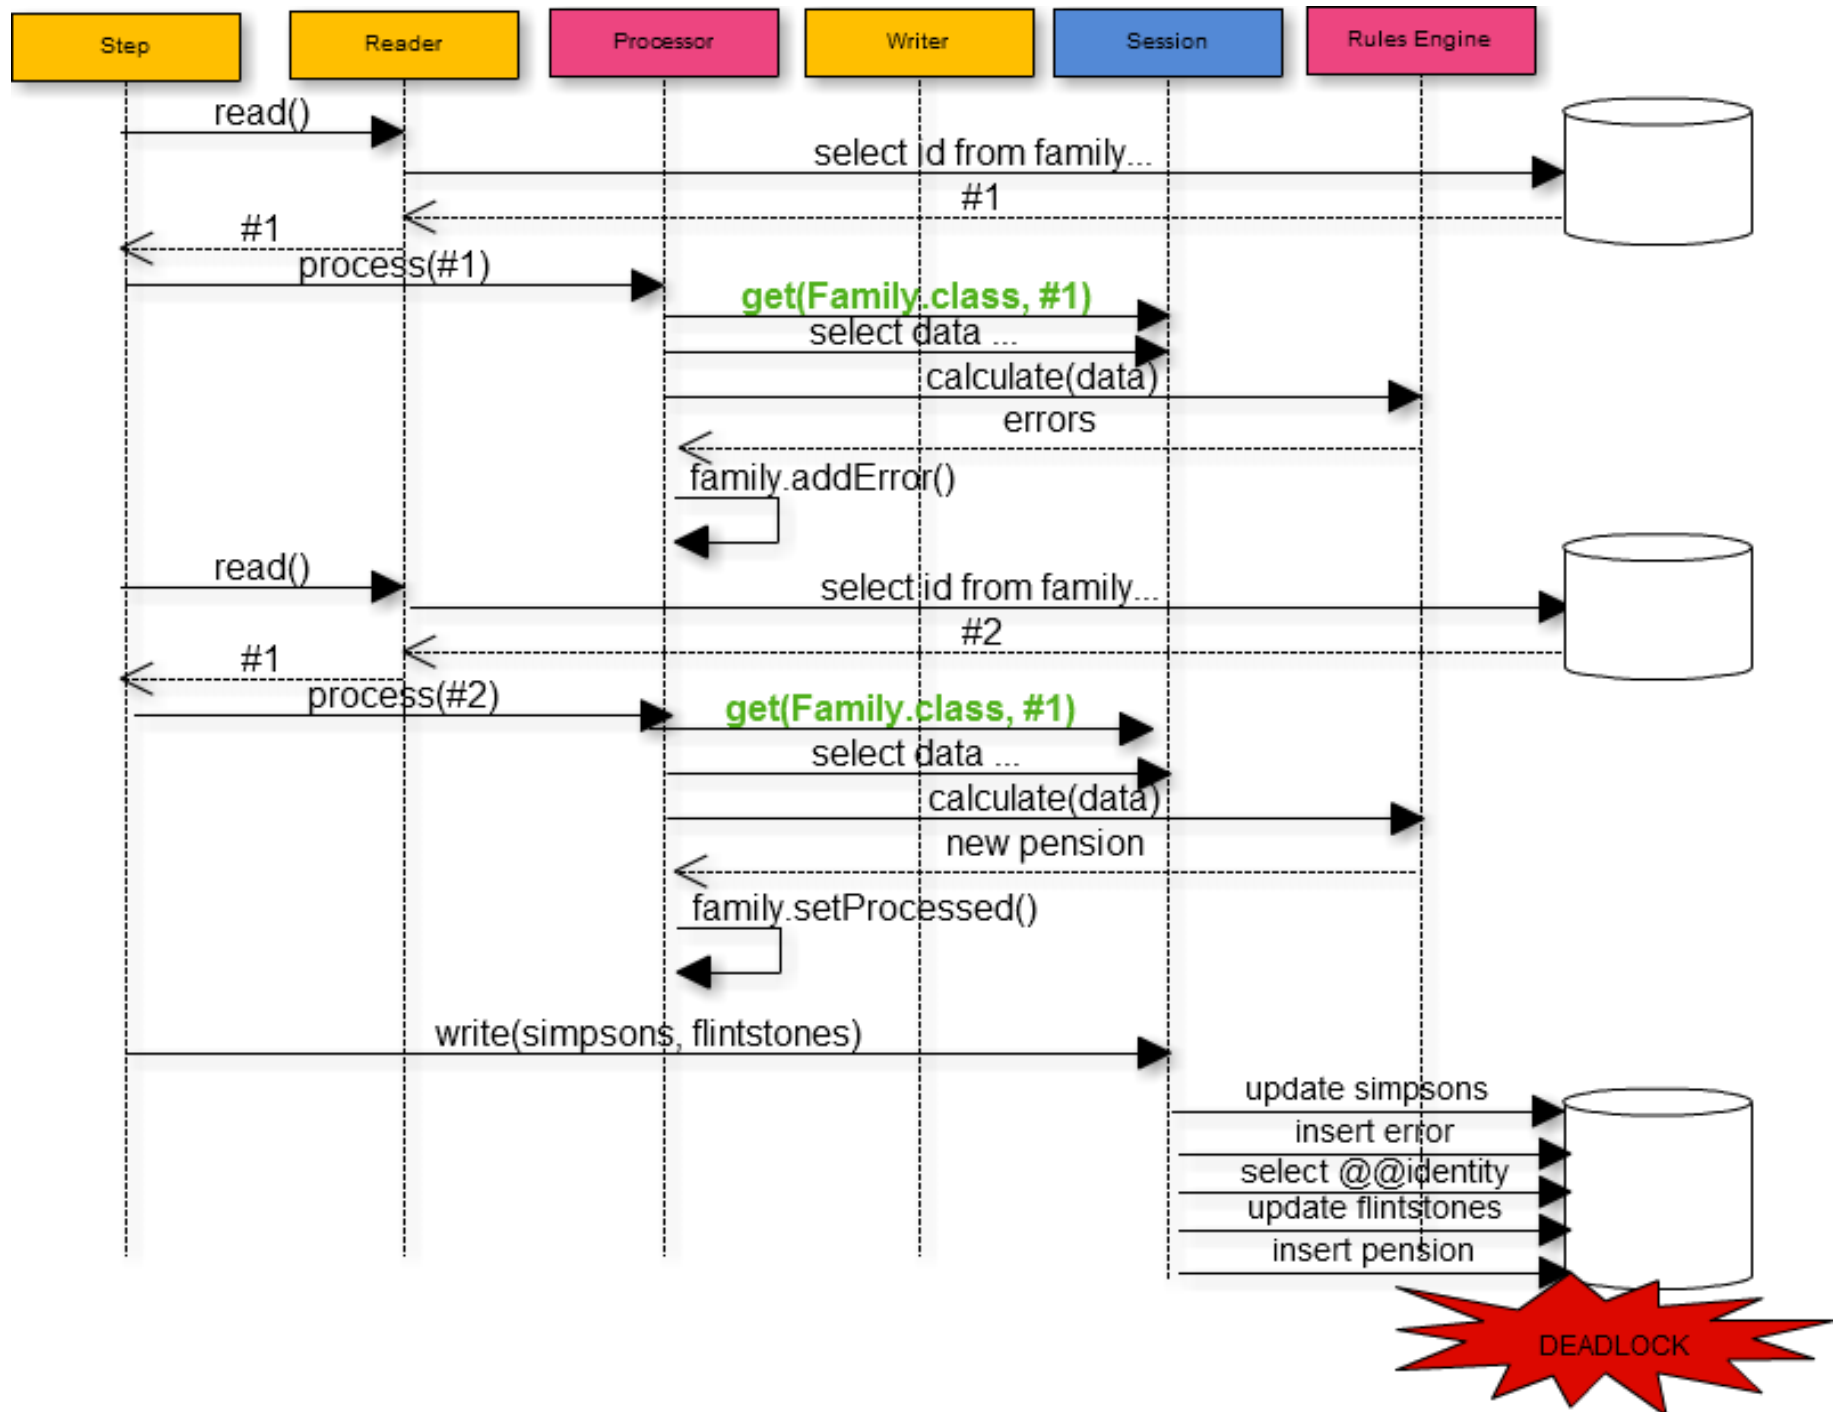


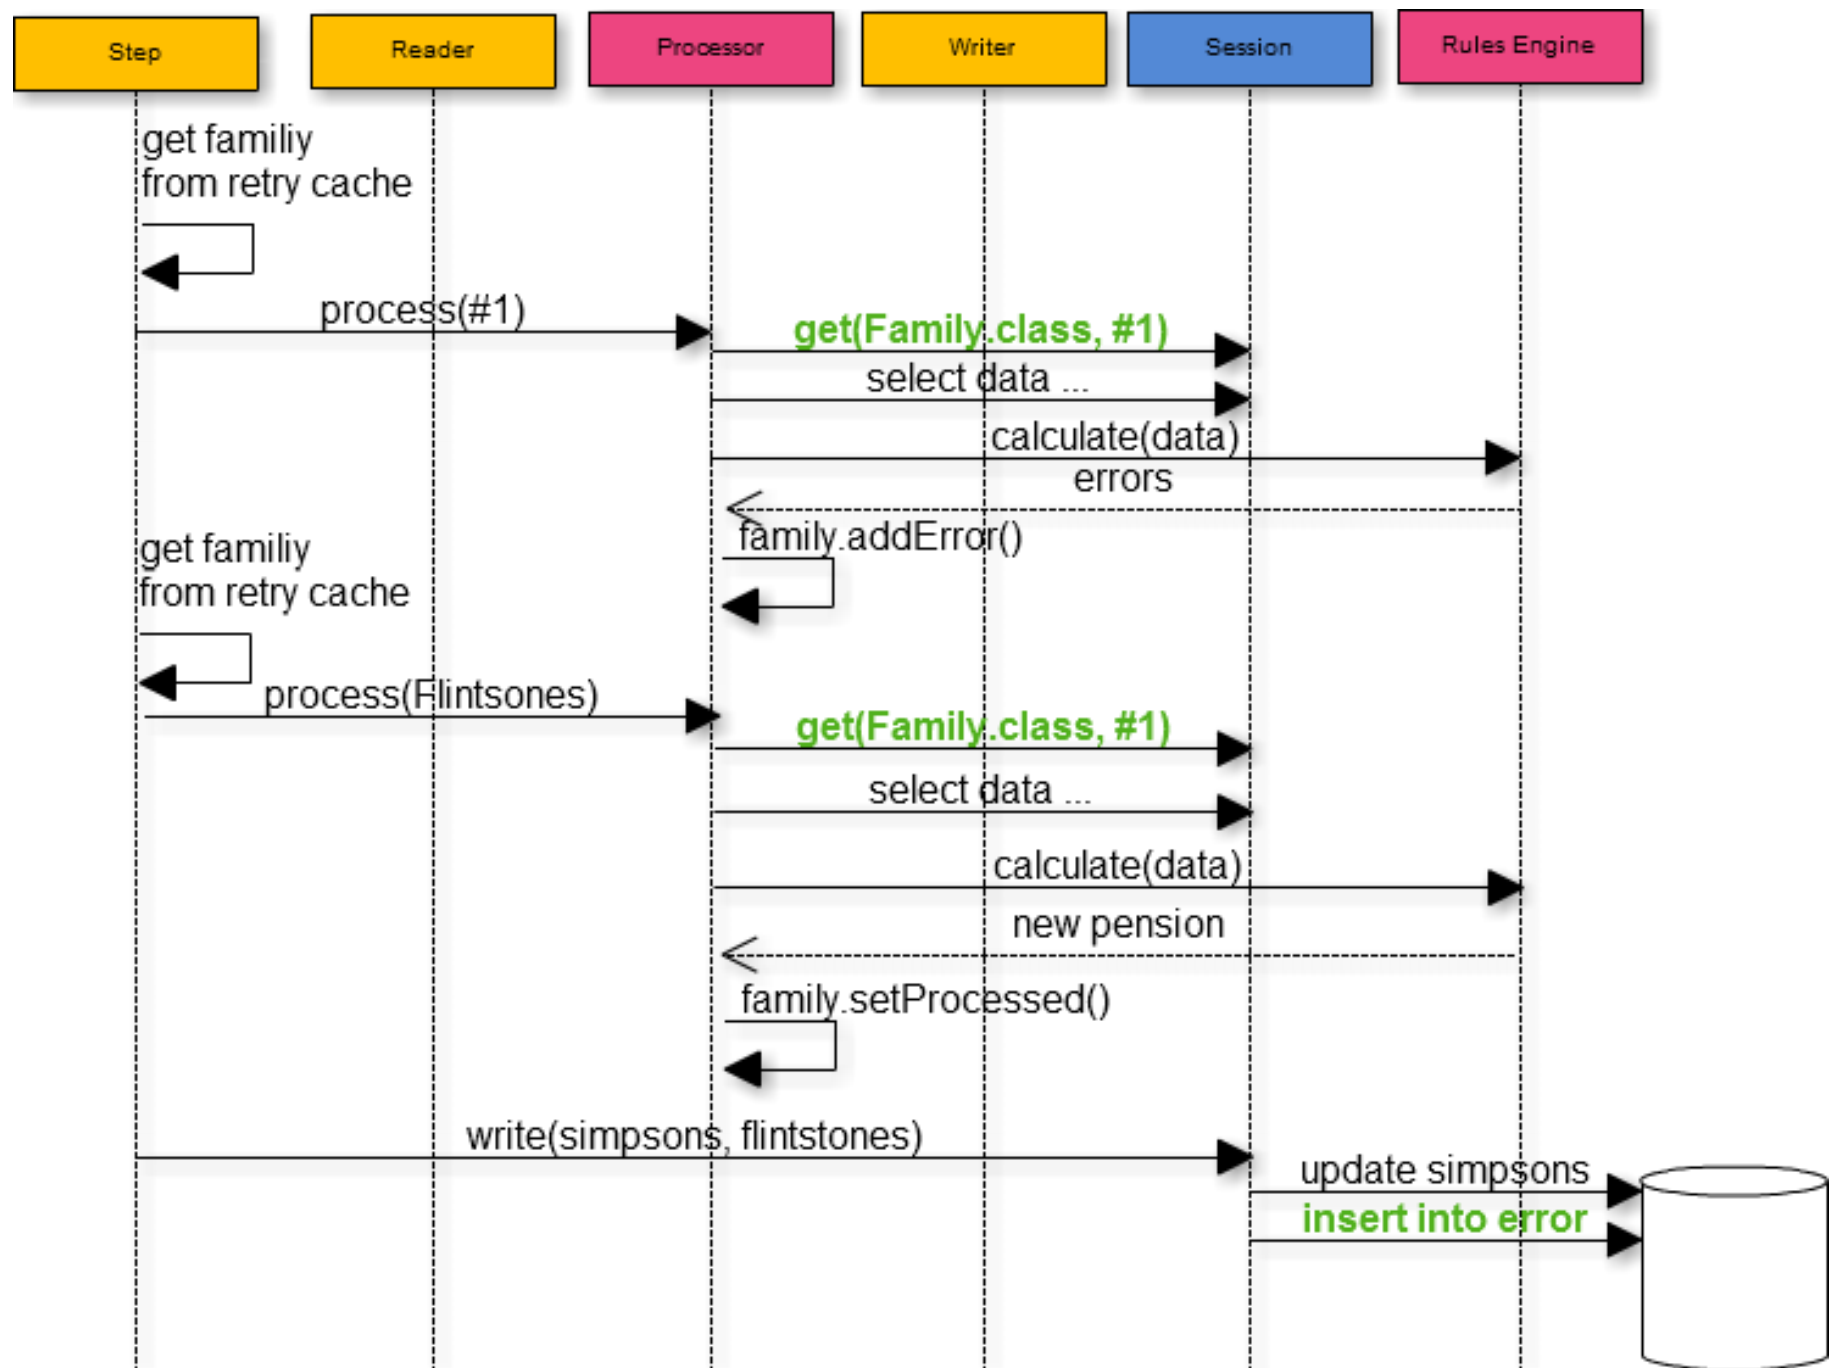
*High performance. Delivered.*

We ditch Hibernate  
...well, almost anyway

# Removing Hibernate from reader

- ItemReader is re-configured to use JDBC
- Fetches primary key from family staging table
- ItemProcessor fetches staging object graph
  - Uses primary key to fetch graph with hibernate
- Primary keys are immutable and stateless







- Performance requirement was 48 hrs
- Completed in 16 hrs
- Used 12 threads
- C-version used 1 thread and ran for 1 week
  - Stopped each morning, started each evening
- A batch that scales with the infrastructure
  - Number of threads is configurable in .properties

# What we would have done differently

- Switched from partitioned to multi-threaded step
  - All work is shared among threads
  - All threads will run until batch completes
  - Avoid idle threads towards the end
  - With partitioning some partitions finished well before others

- Do not use Hibernate in the ItemReader
- Test parallelization early
- Monitor your SQLs
  - Frequently called
  - Long running
- Become friends with your DBA
- There is no reason to let Java be the bottle neck
  - Increase thread count until DB becomes the bottle neck



@mortenag



[www.github.com/magott](https://www.github.com/magott)



[www.andersen-gott.com](http://www.andersen-gott.com)

- If one lazily loaded entity is fetched, they are all fetched – in one query

```
@Entity
public class Order {

    @BatchSize(size=20)
    private Collection<Item> items;
```